

# Operating Systems

Associate Prof. Yongkun Li

中科大-计算机学院 副教授

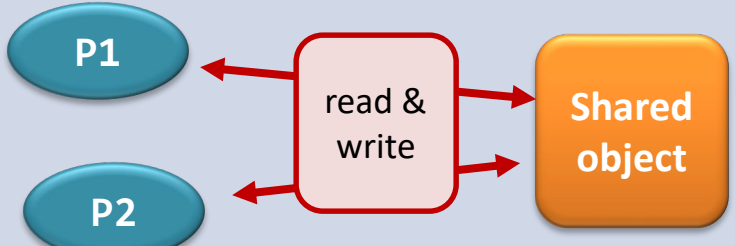
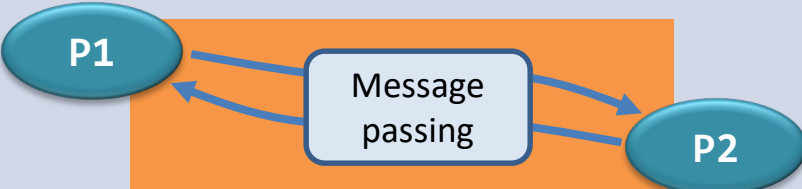
<http://staff.ustc.edu.cn/~ykli>

Ch5


Process Communication & Synchronization

-Part 2

# Summary on IPC models – another point of view


Shared Objects	Message Passing
 <p>The diagram shows two processes, P1 and P2, represented as blue ovals on the left. To their right is a red rounded rectangle labeled 'read &amp; write' and an orange rounded rectangle labeled 'Shared object'. Red arrows point from both P1 and P2 to the 'read &amp; write' box, and from this box to the 'Shared object' box, indicating that both processes interact with a common shared resource.</p>	 <p>The diagram shows two processes, P1 and P2, represented as blue ovals. Between them is a light blue rounded rectangle labeled 'Message passing' set against an orange background. Blue arrows show a bidirectional flow of communication between P1 and P2 through the message passing mechanism.</p>
<p><b><u>Challenge.</u></b> Coordination can only be done by detecting the status of the shared object. <i>E.g., is the pipe empty / full?</i></p>	<p><b><u>Challenge.</u></b> Coordination relies on the reliability and the efficiency of the communication medium (and protocol).</p>
<p>E.g., pipes, shared memory, and regular files.</p>	<p>E.g., socket programming, message passing interface (MPI) library.</p>

# IPC problem: Race condition



# Evil source: the shared objects

- Pipe is implemented with the thought that **there may be more than one process accessing it “at the same time”**
- For shared memory and files, **concurrent access may yield unpredictable outcomes**



# Understanding the problem...


## High-level language for Program A

```
1 attach to the shared memory X;  
2 add 10 to X;  
3 exit;
```

## High-level language for Program B

```
1 attach to the shared memory X;  
2 minus 10 to X;  
3 exit;
```

## The Scenario



Guess what the final result should be?

It may be 10, 0 or 20, can you believe it?

# Understanding the problem...


## High-level language for Program A

```
1 attach to the shared memory X;  
2 add 10 to X;  
3 exit;
```

## High-level language for Program B

```
1 attach to the shared memory X;  
2 minus 10 to X;  
3 exit;
```

## The Scenario



Remember the flow of executing a program and the system hierarchy?

# Understanding the problem...

## High-level language for Program A


```
1 attach to the shared memory X;  
2 add 10 to X;  
3 exit;
```

This operation  
is not atomic

## Partial low-level language for Program A

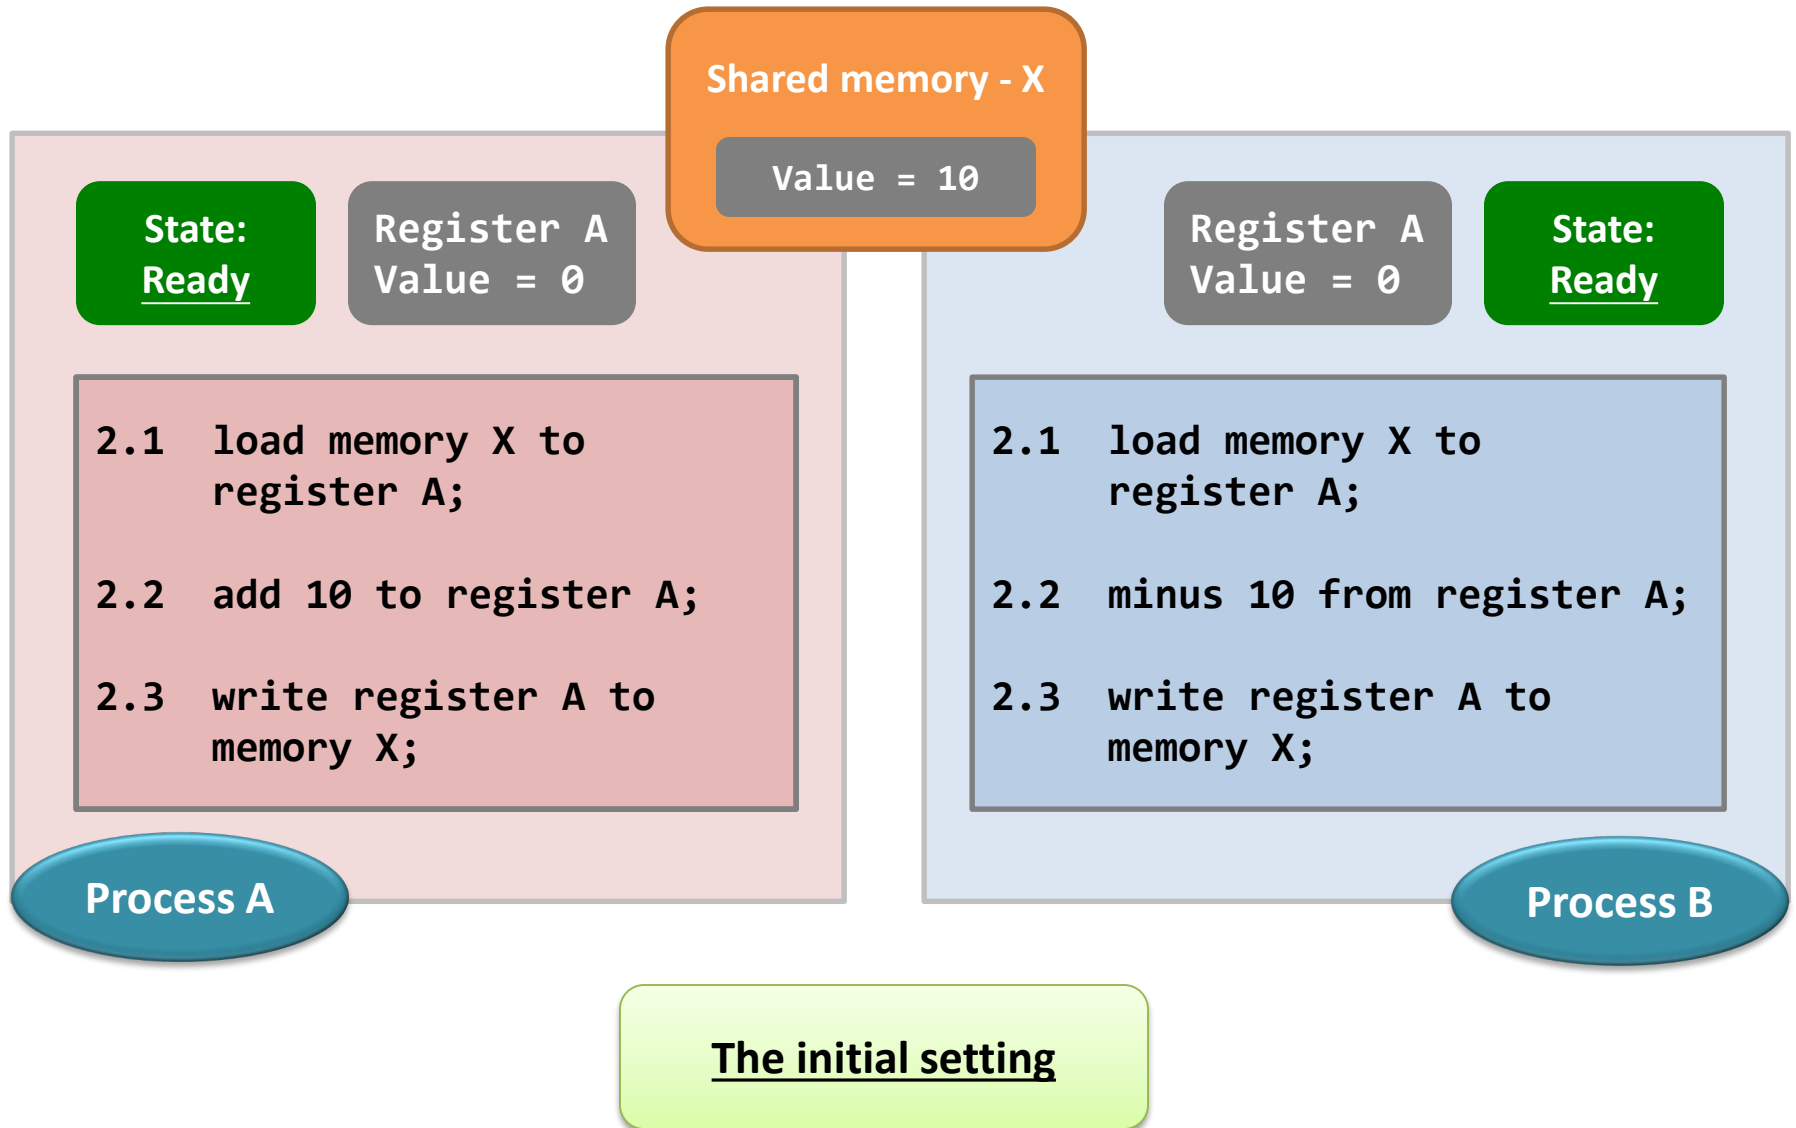
```
1 attach to the shared memory X;  
.....  
2.1 load memory X to register A;  
2.2 add 10 to register A;  
2.3 write register A to memory X;  
.....  
3 exit;
```

## The Scenario



Guess what? This code block is evil!

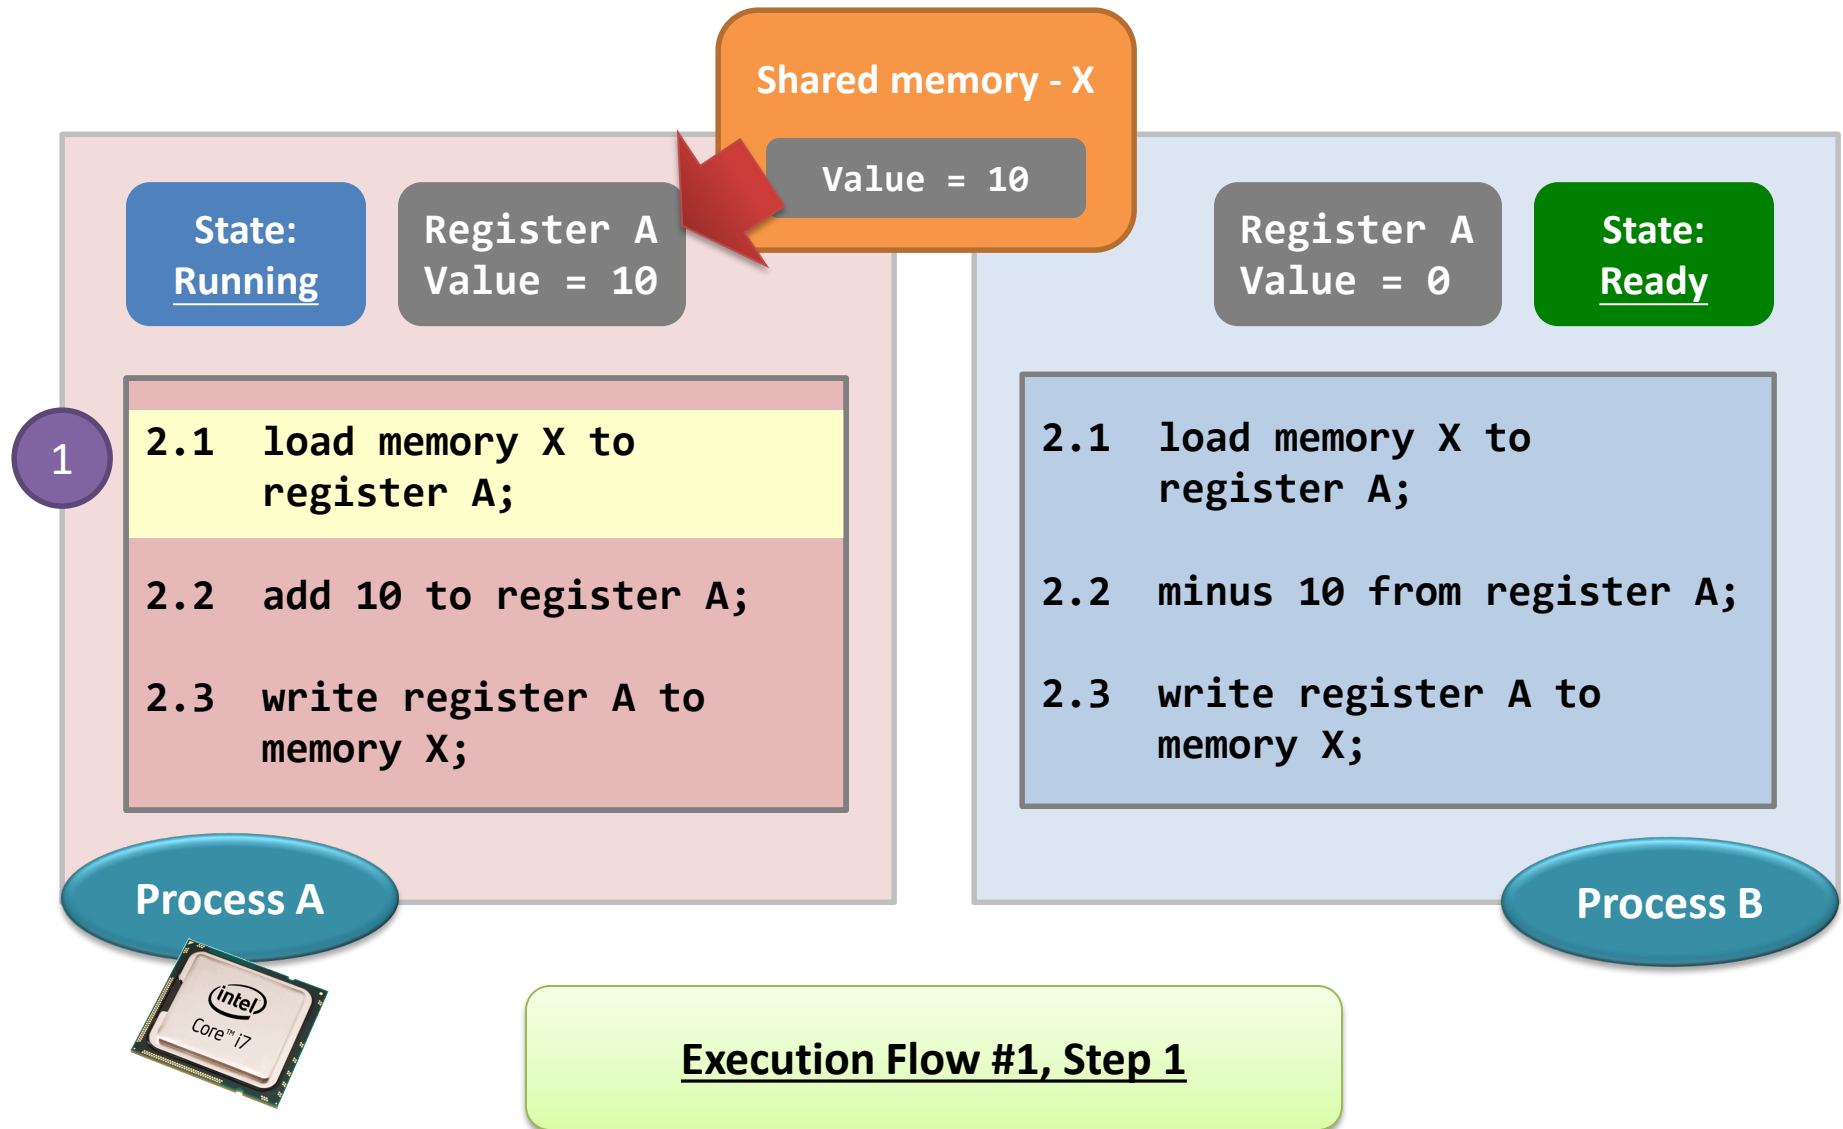
# Understanding the problem...



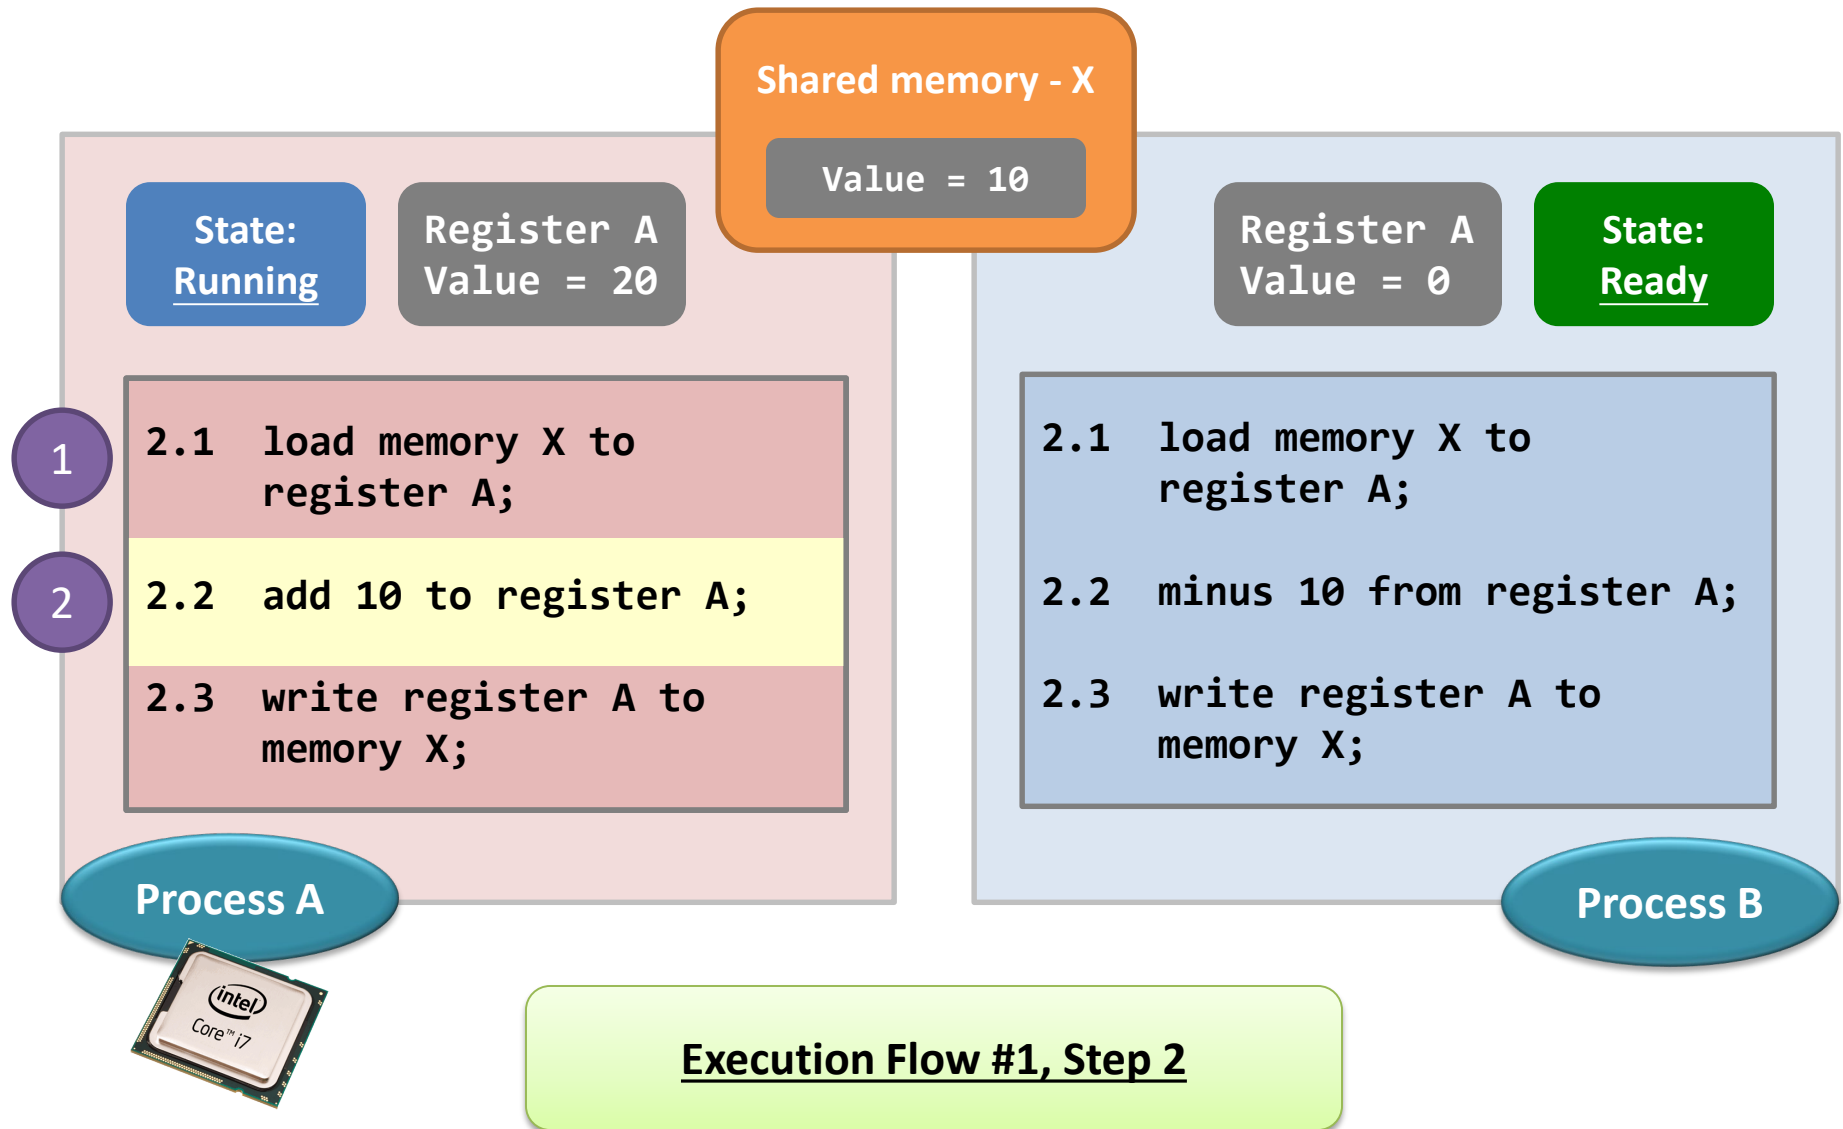


## **Execution Flow #1**

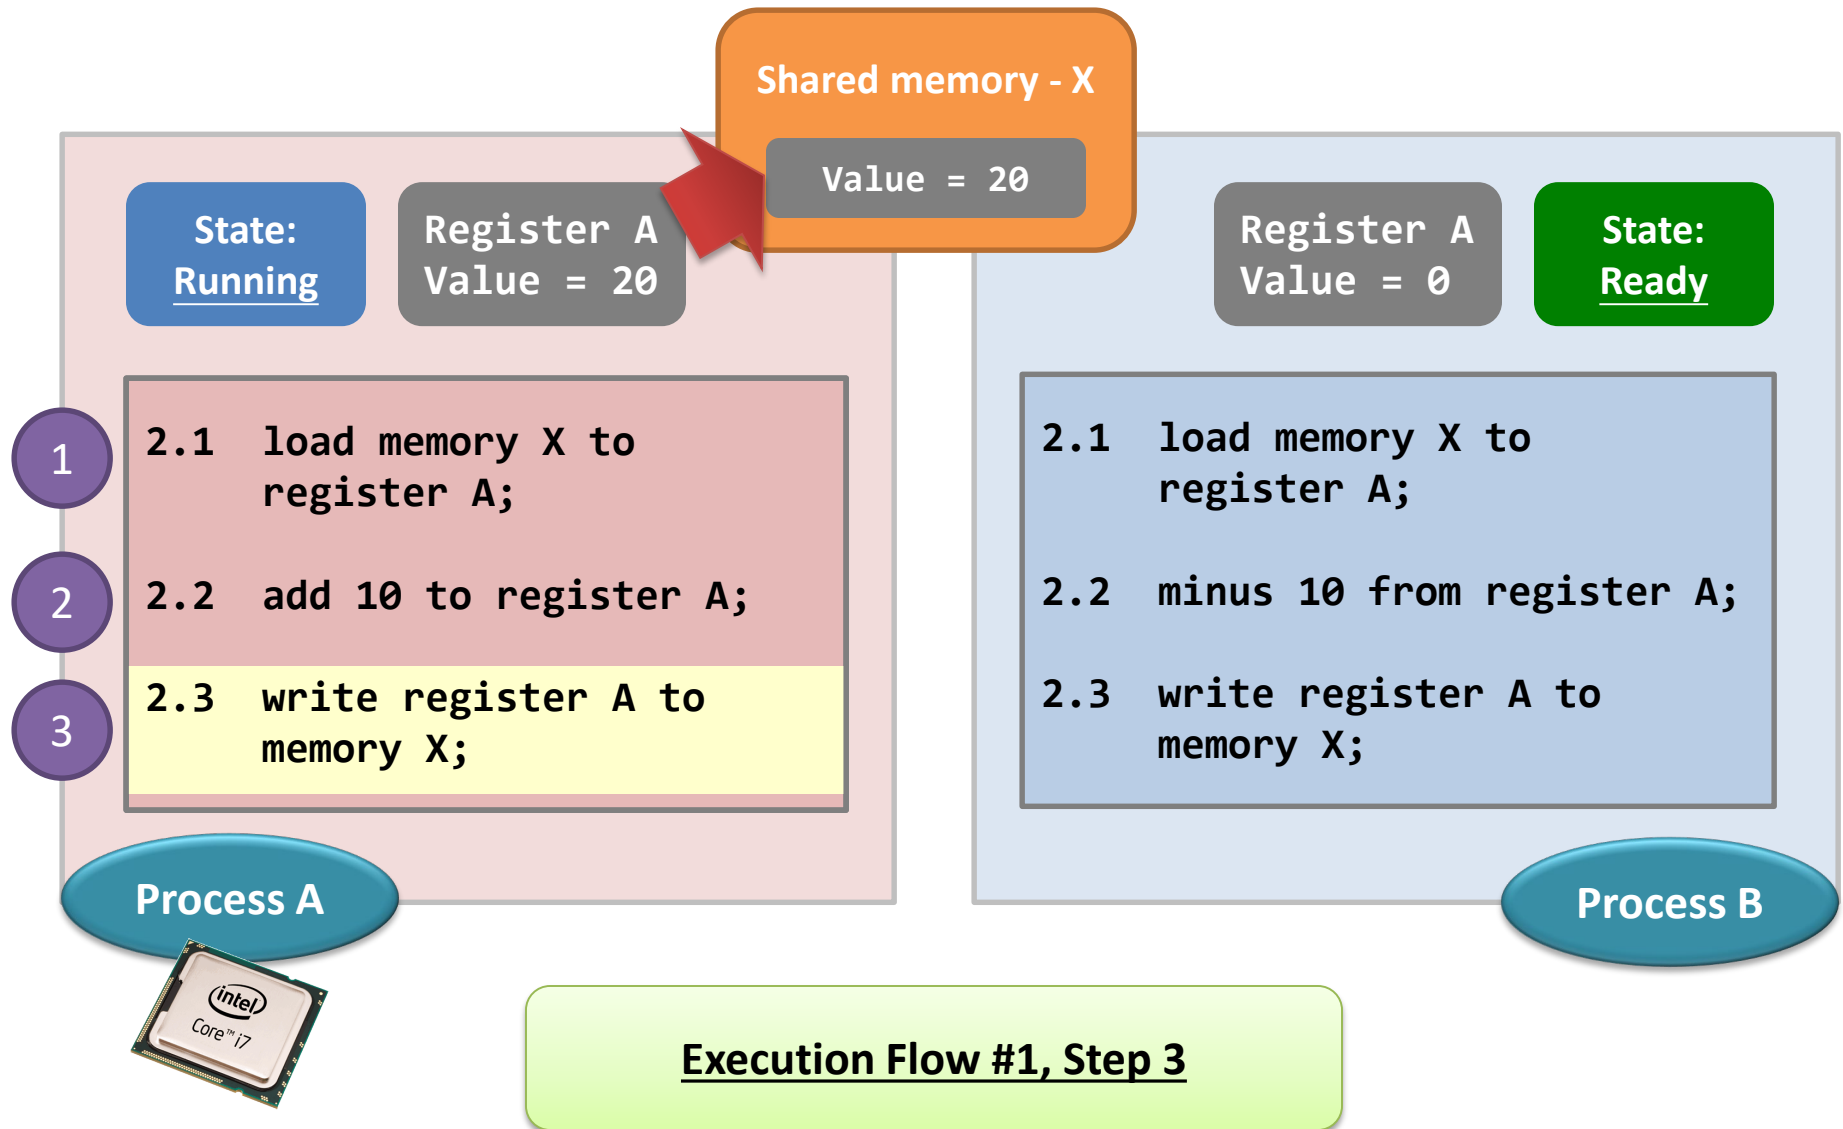
# Problem not yet arise...



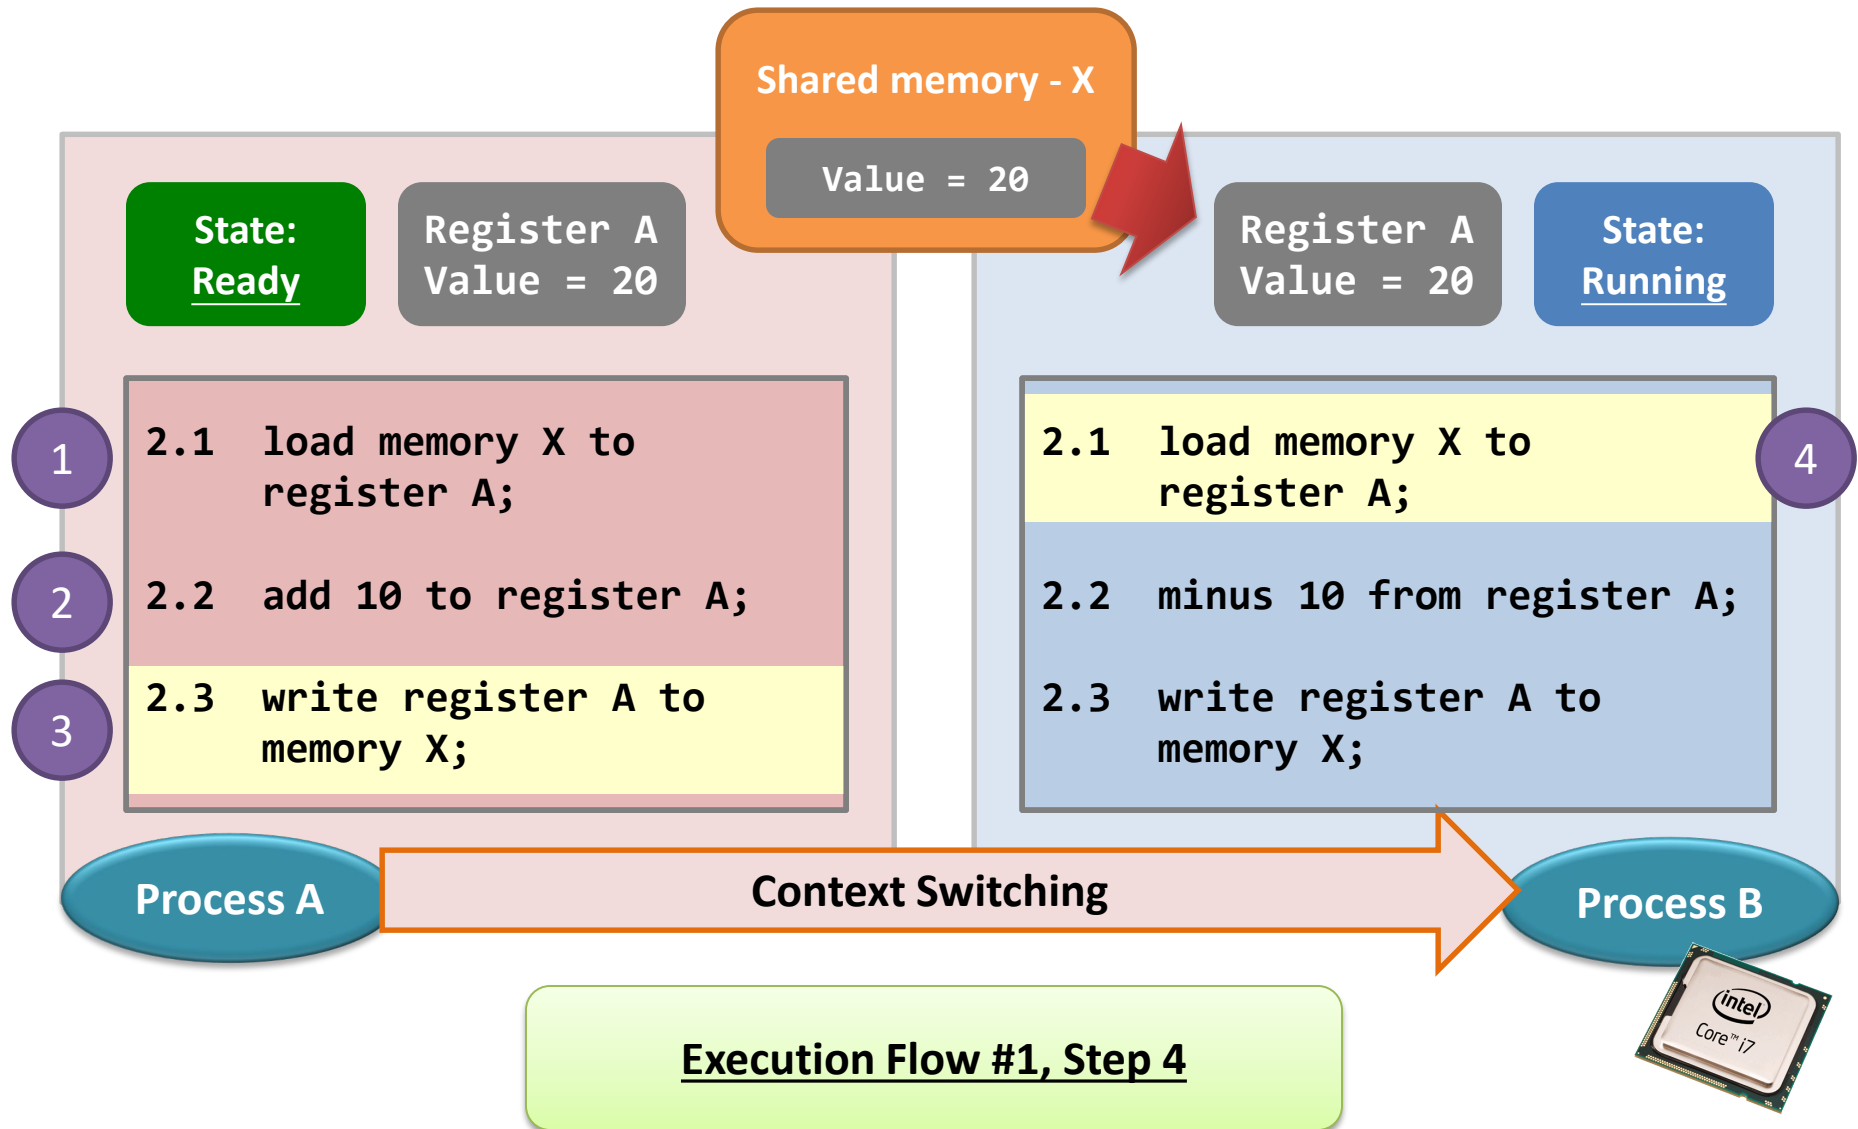
# Problem not yet arise...



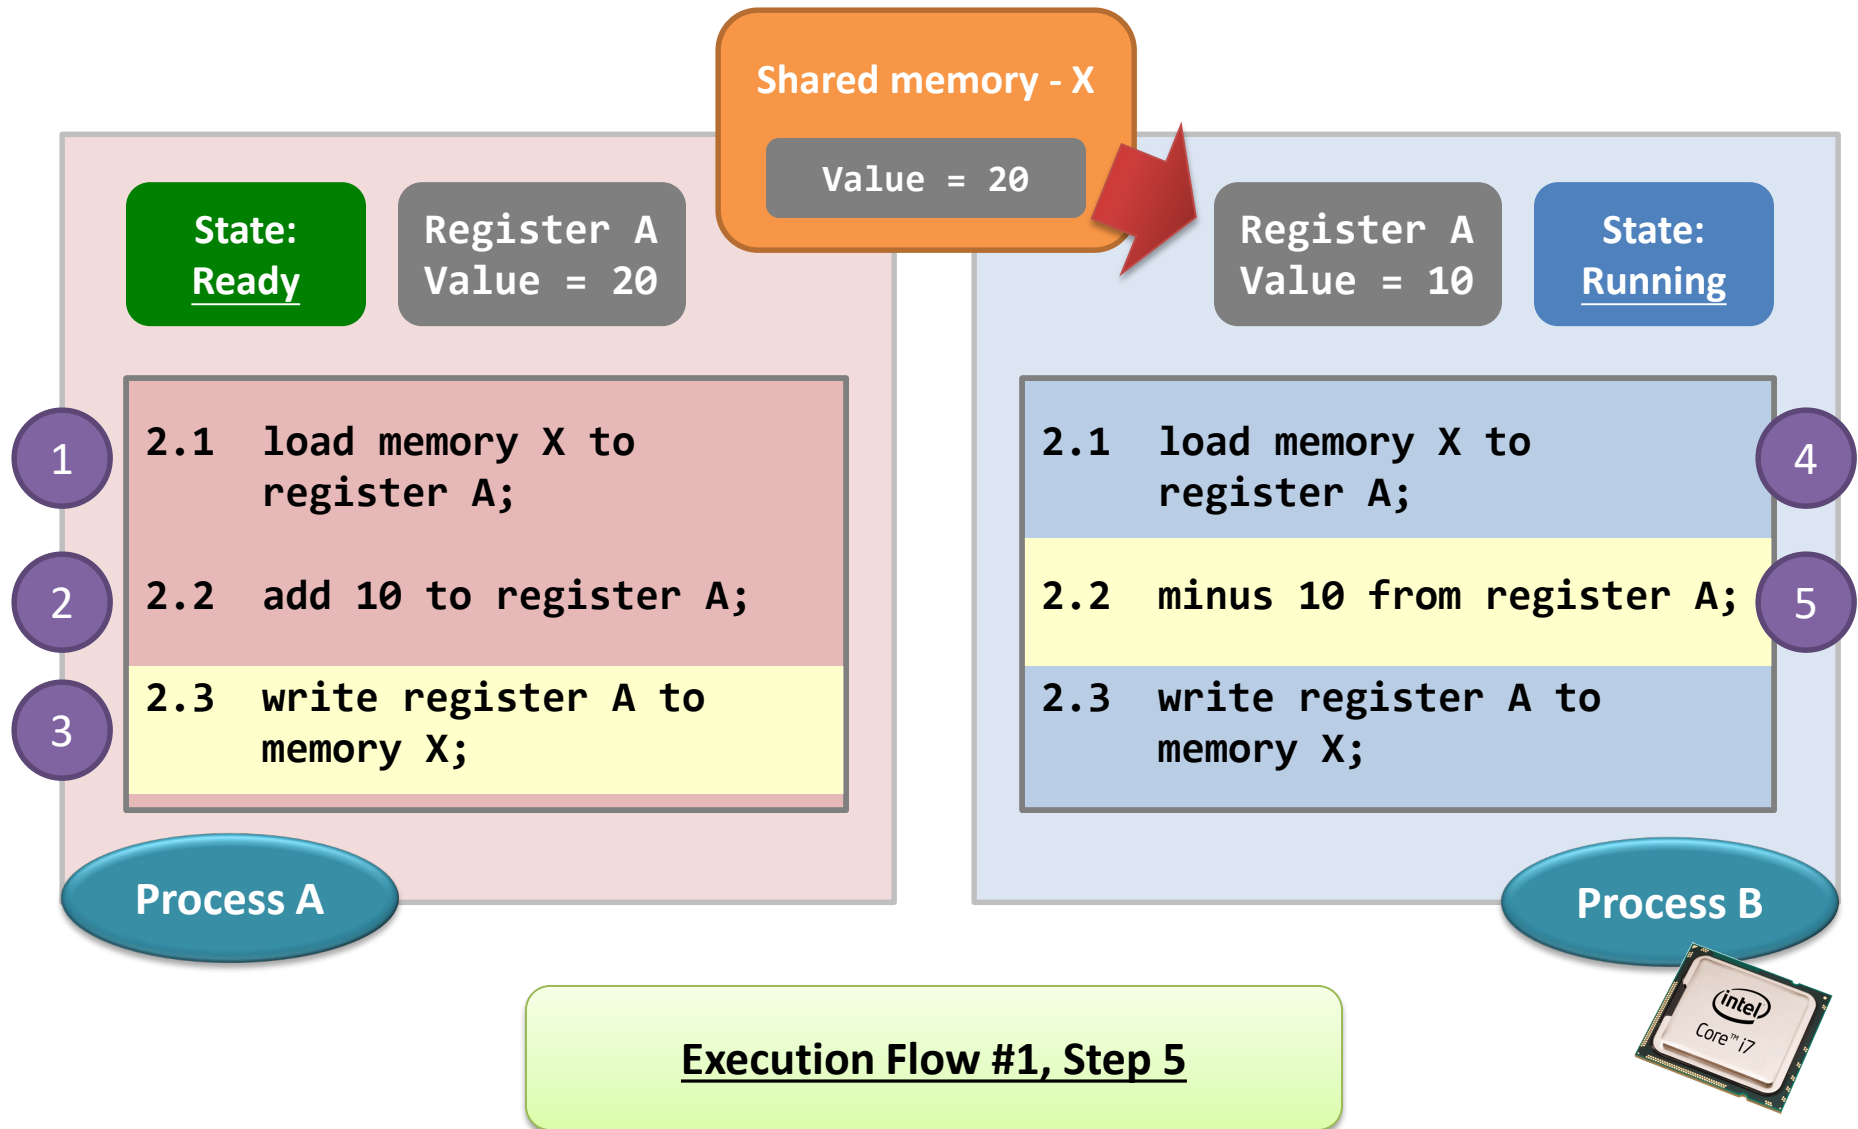
# Problem not yet arise...



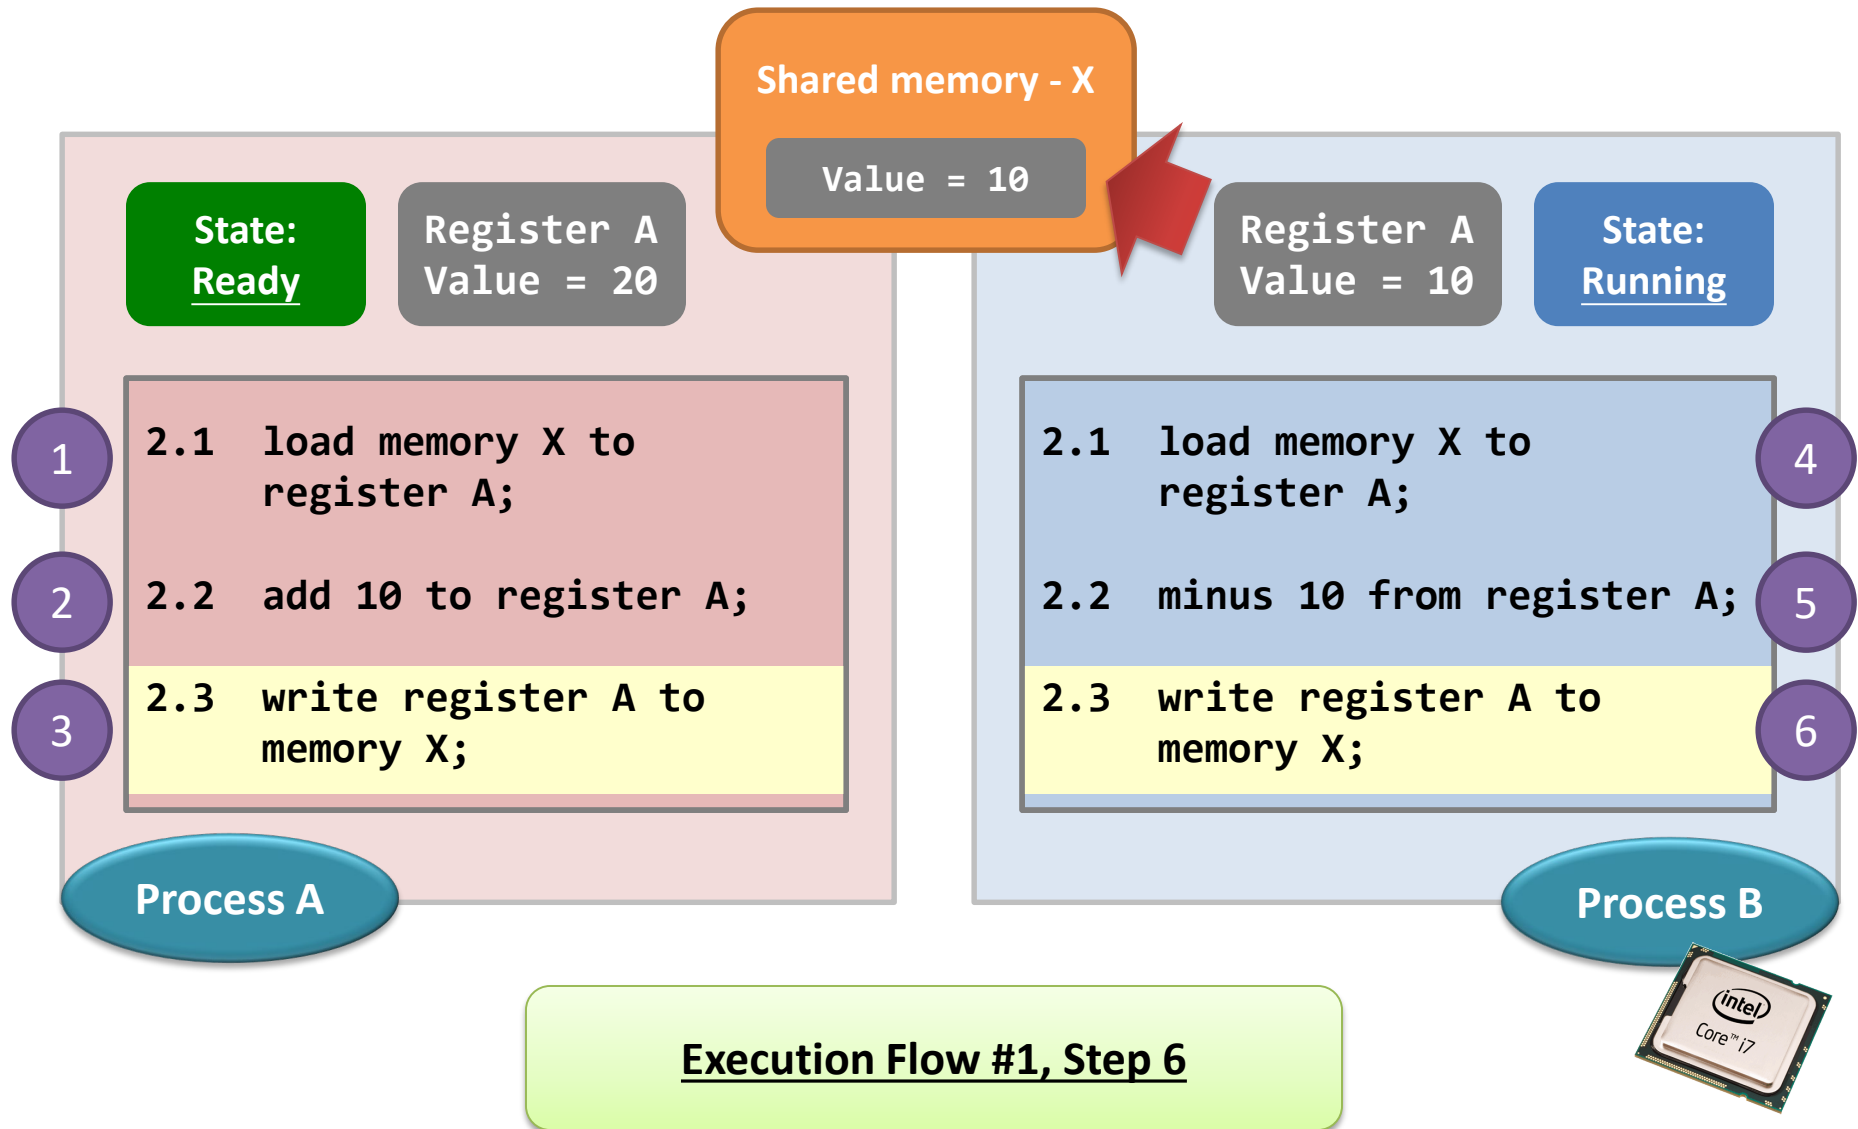
# Problem not yet arise...



# Problem not yet arise...



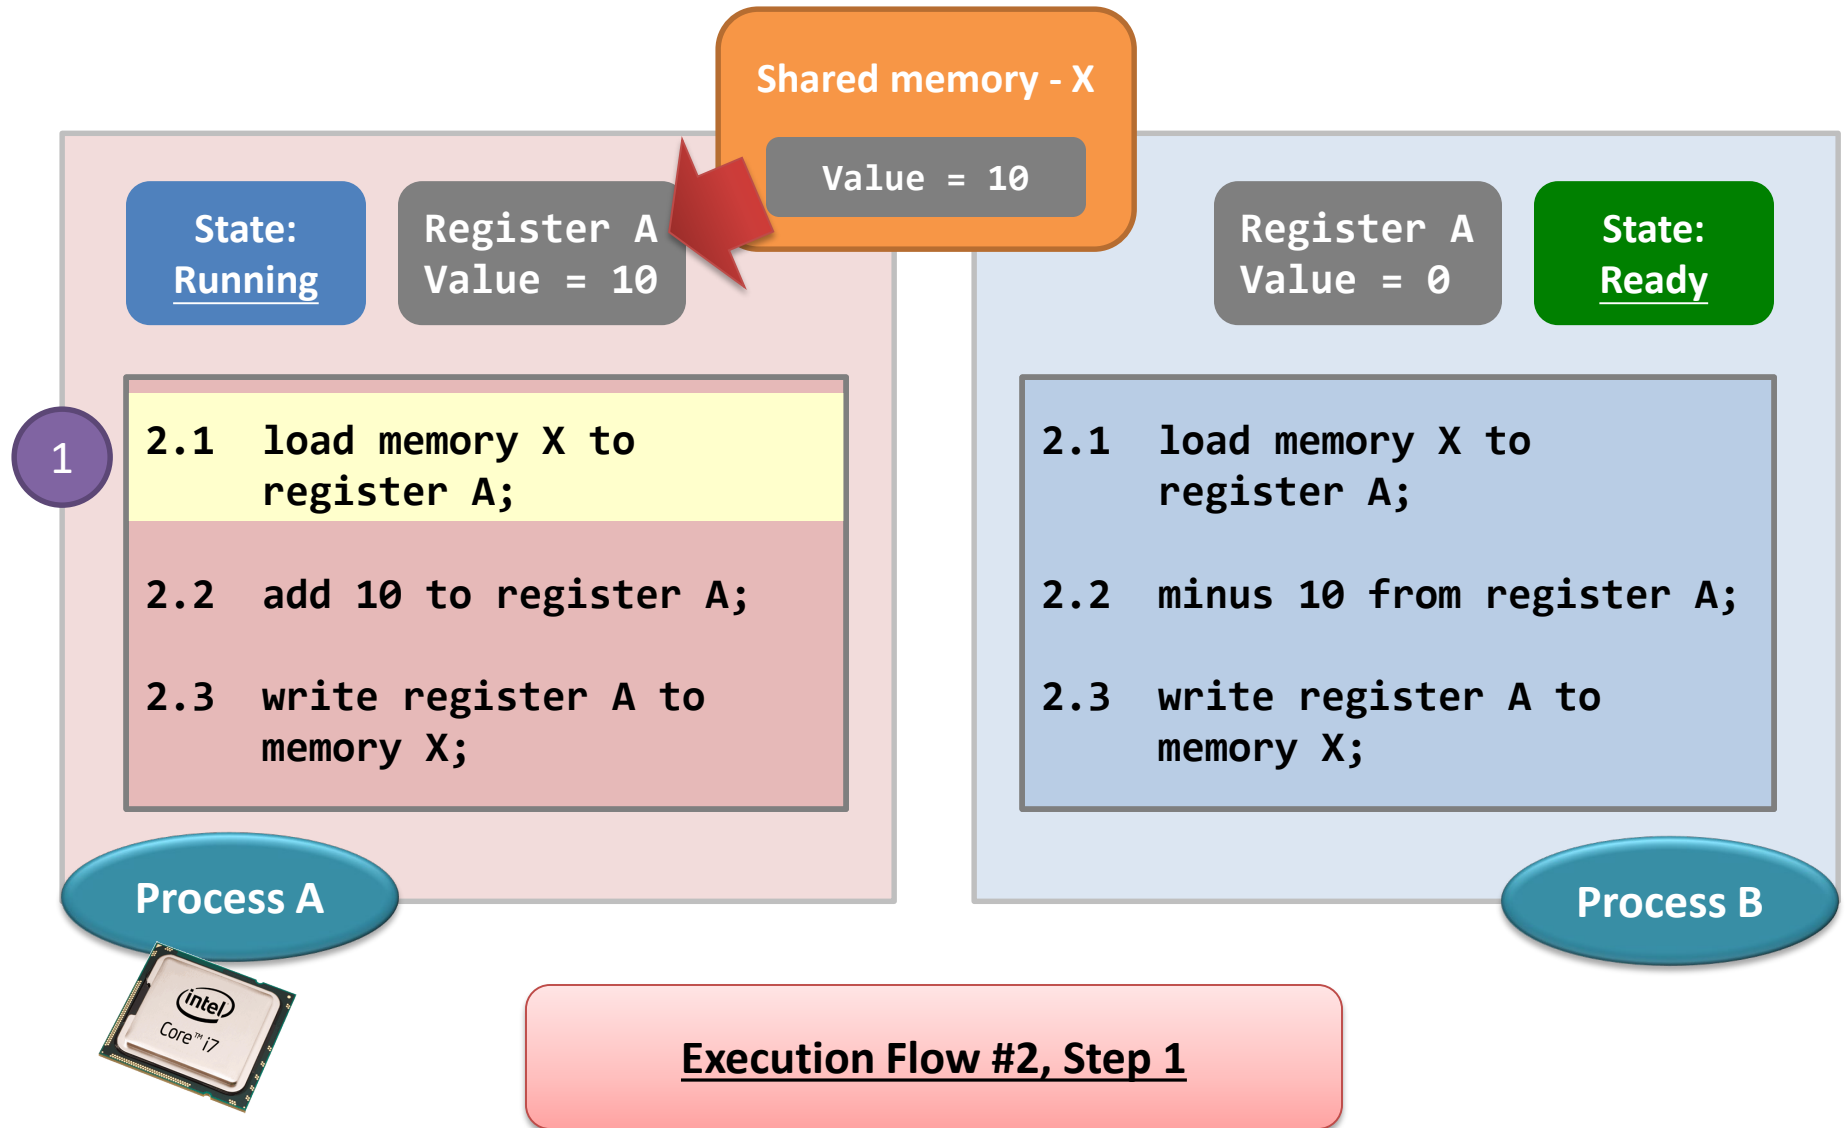
# Problem not yet arise...



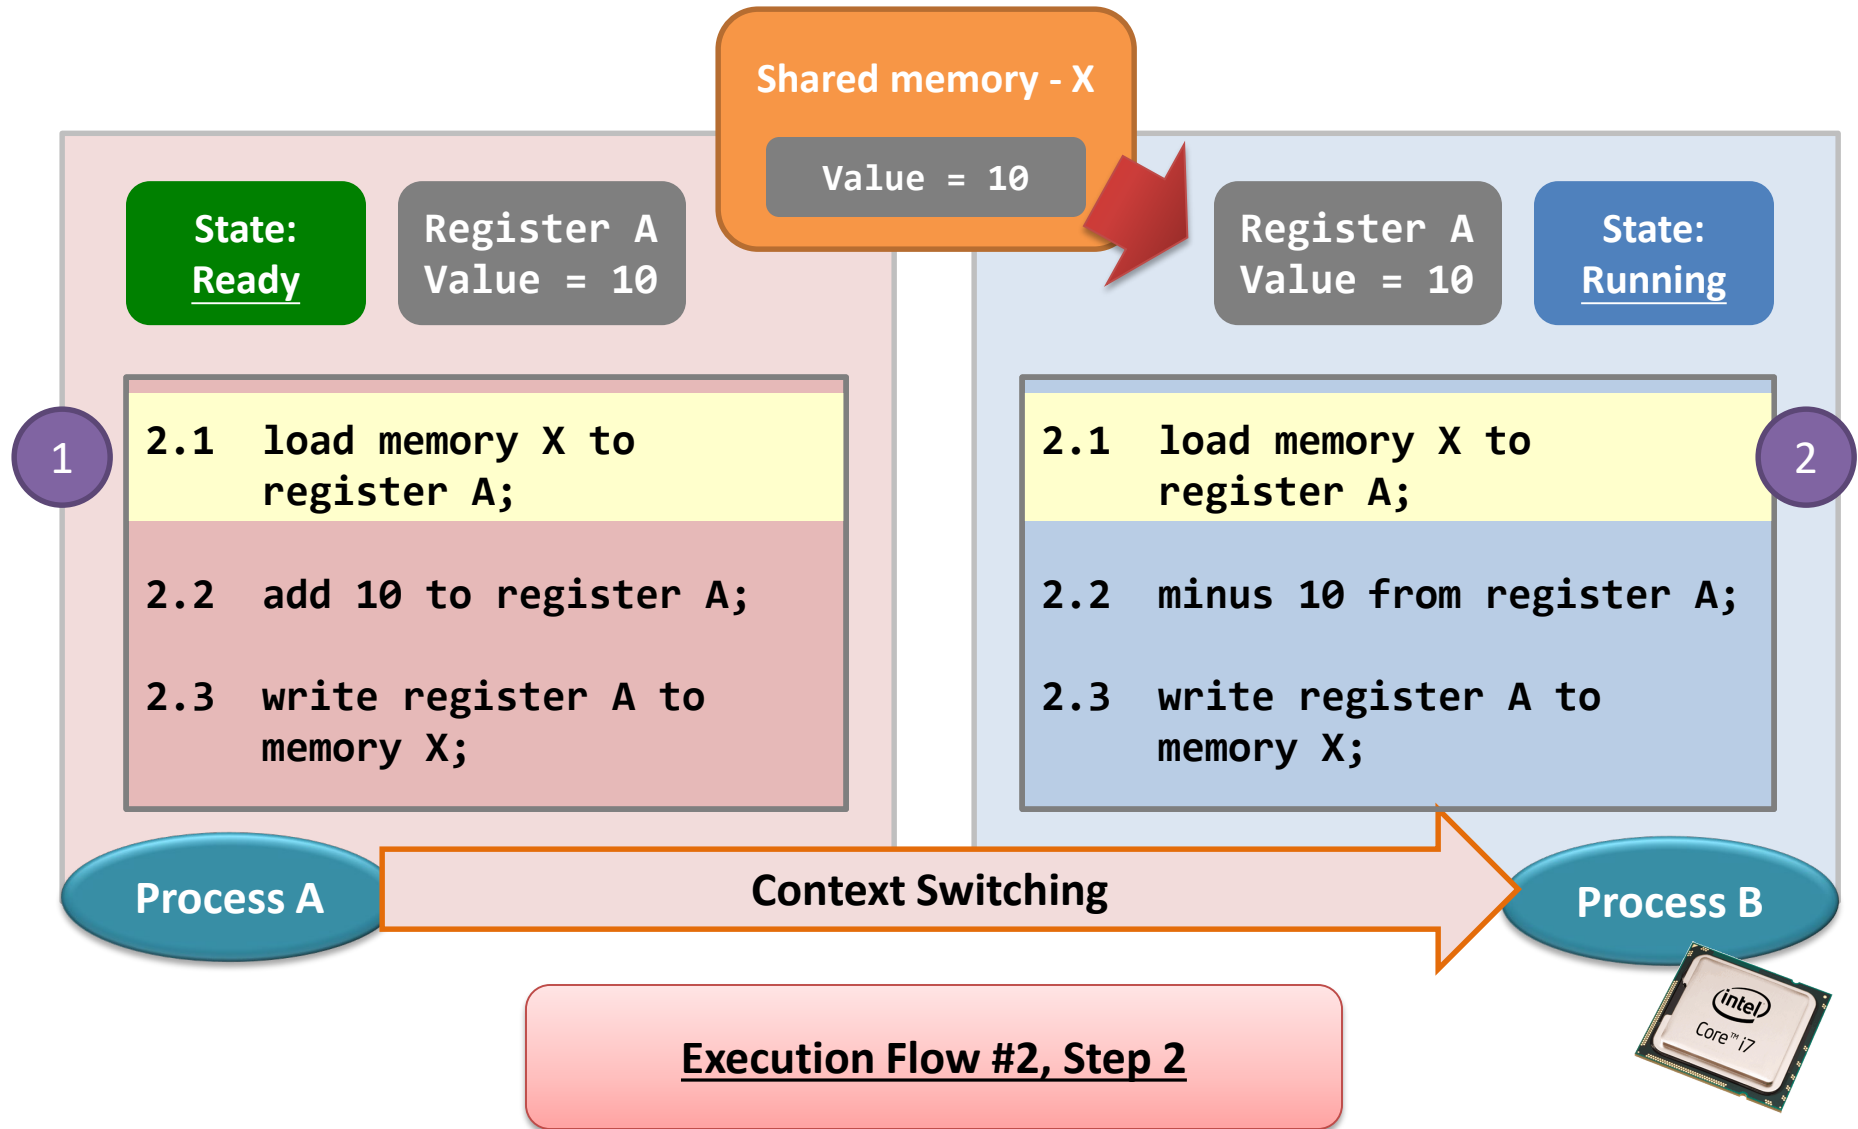
## **Execution Flow #2**



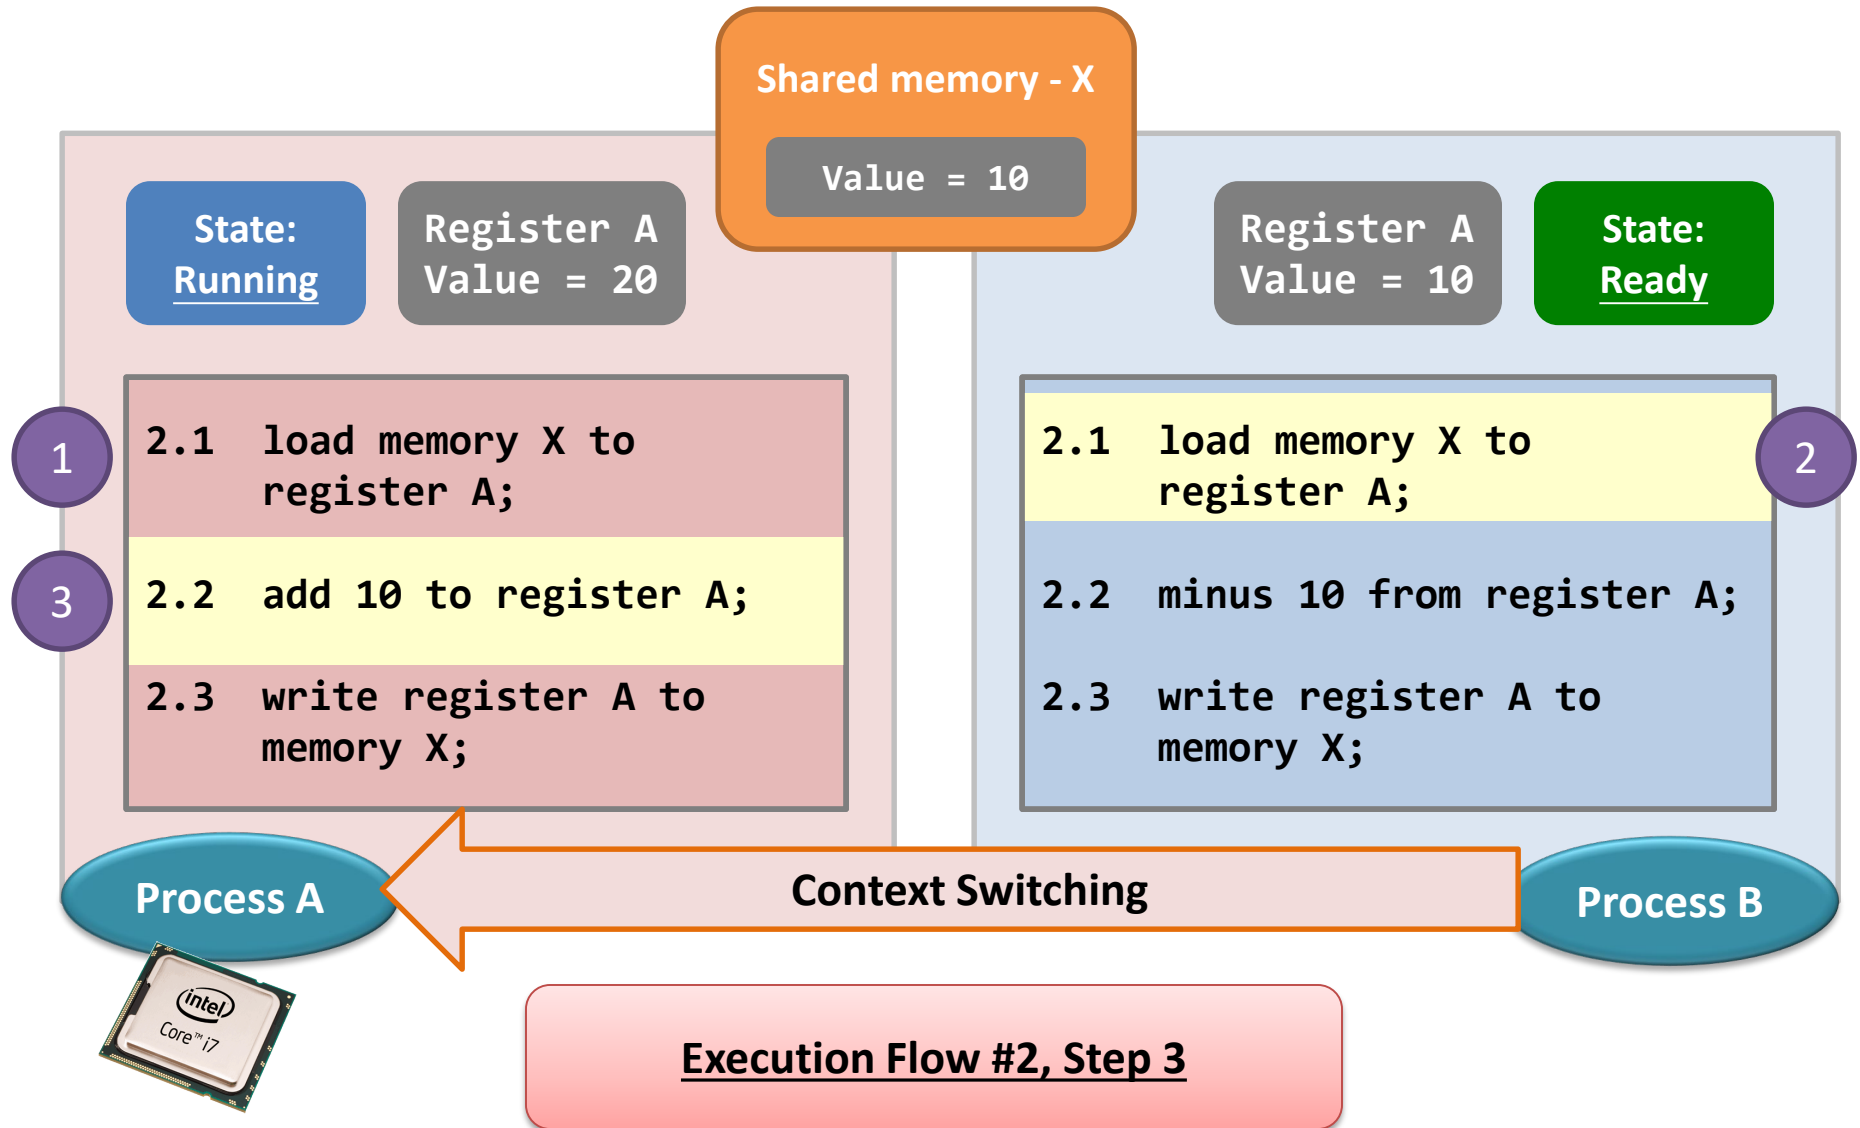
# Problem arise...



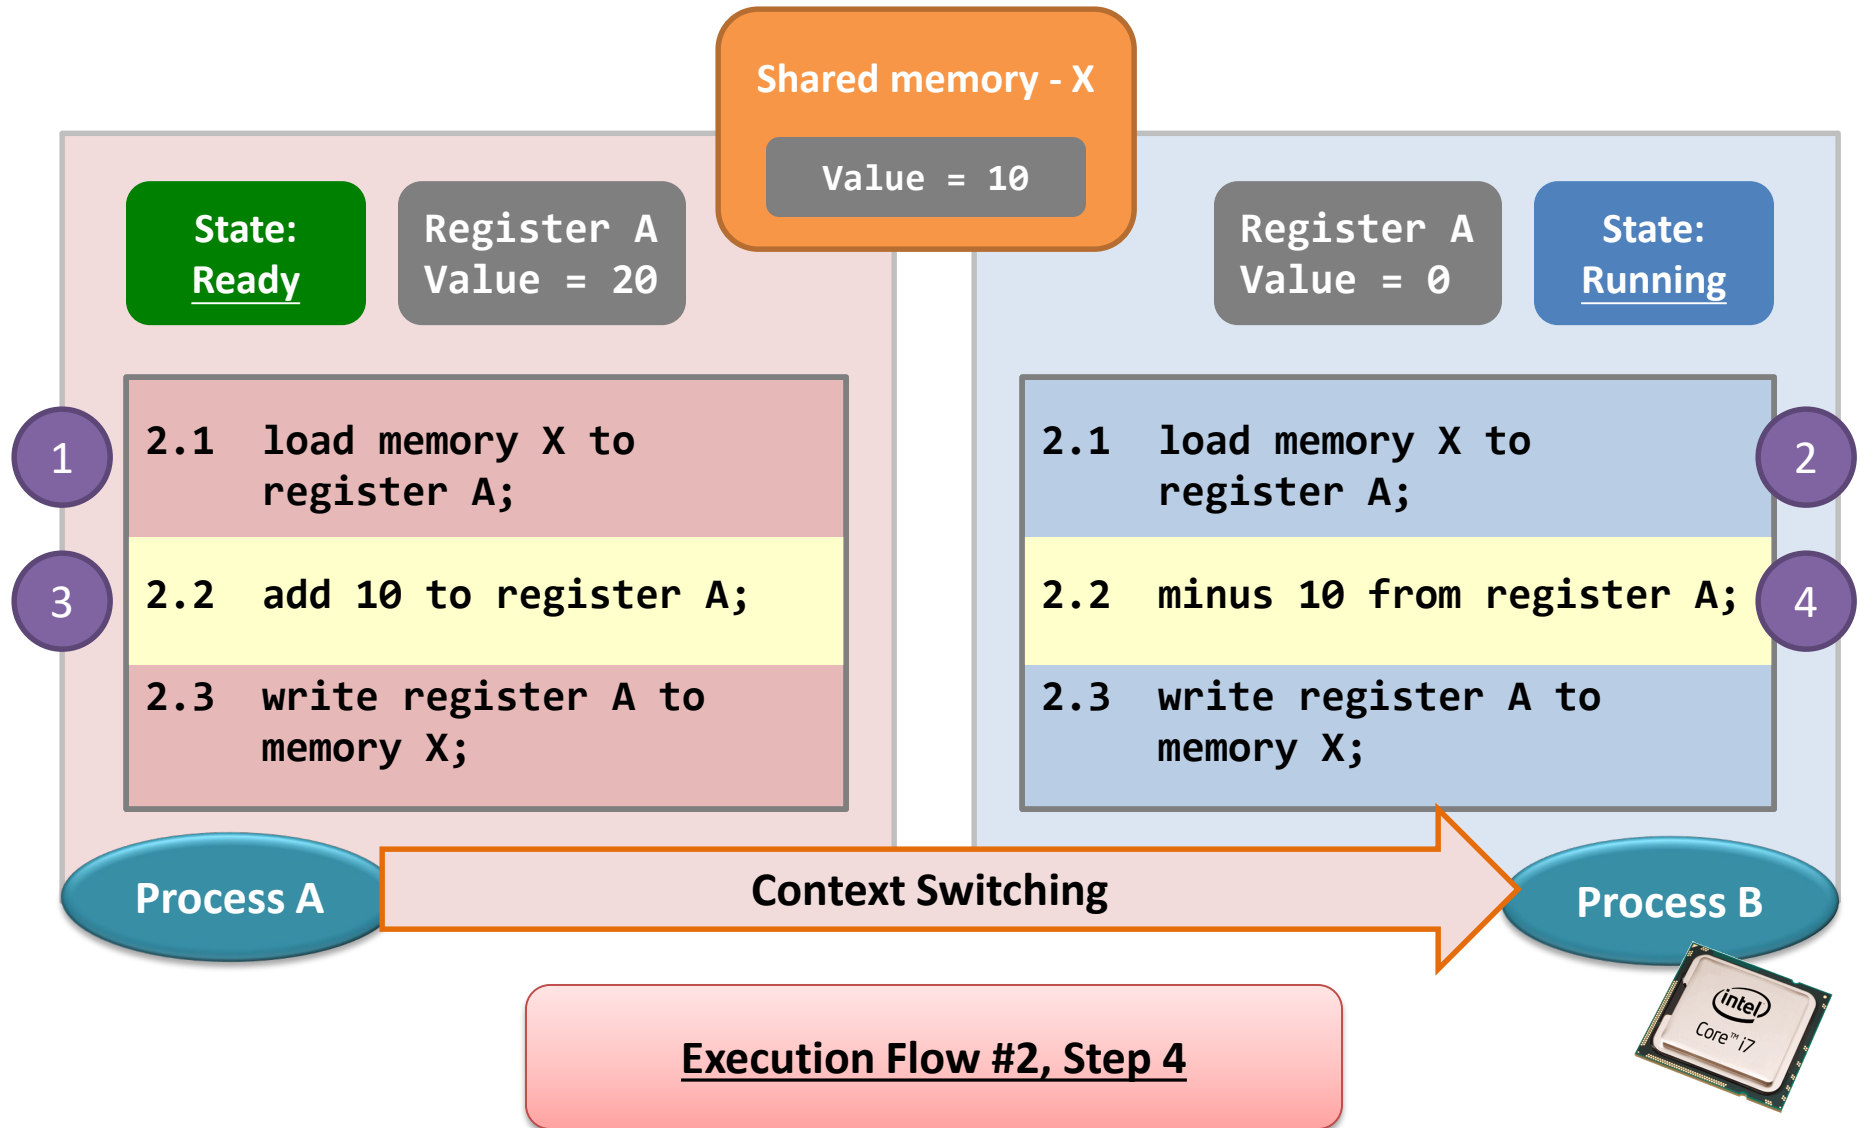
# Problem arise...



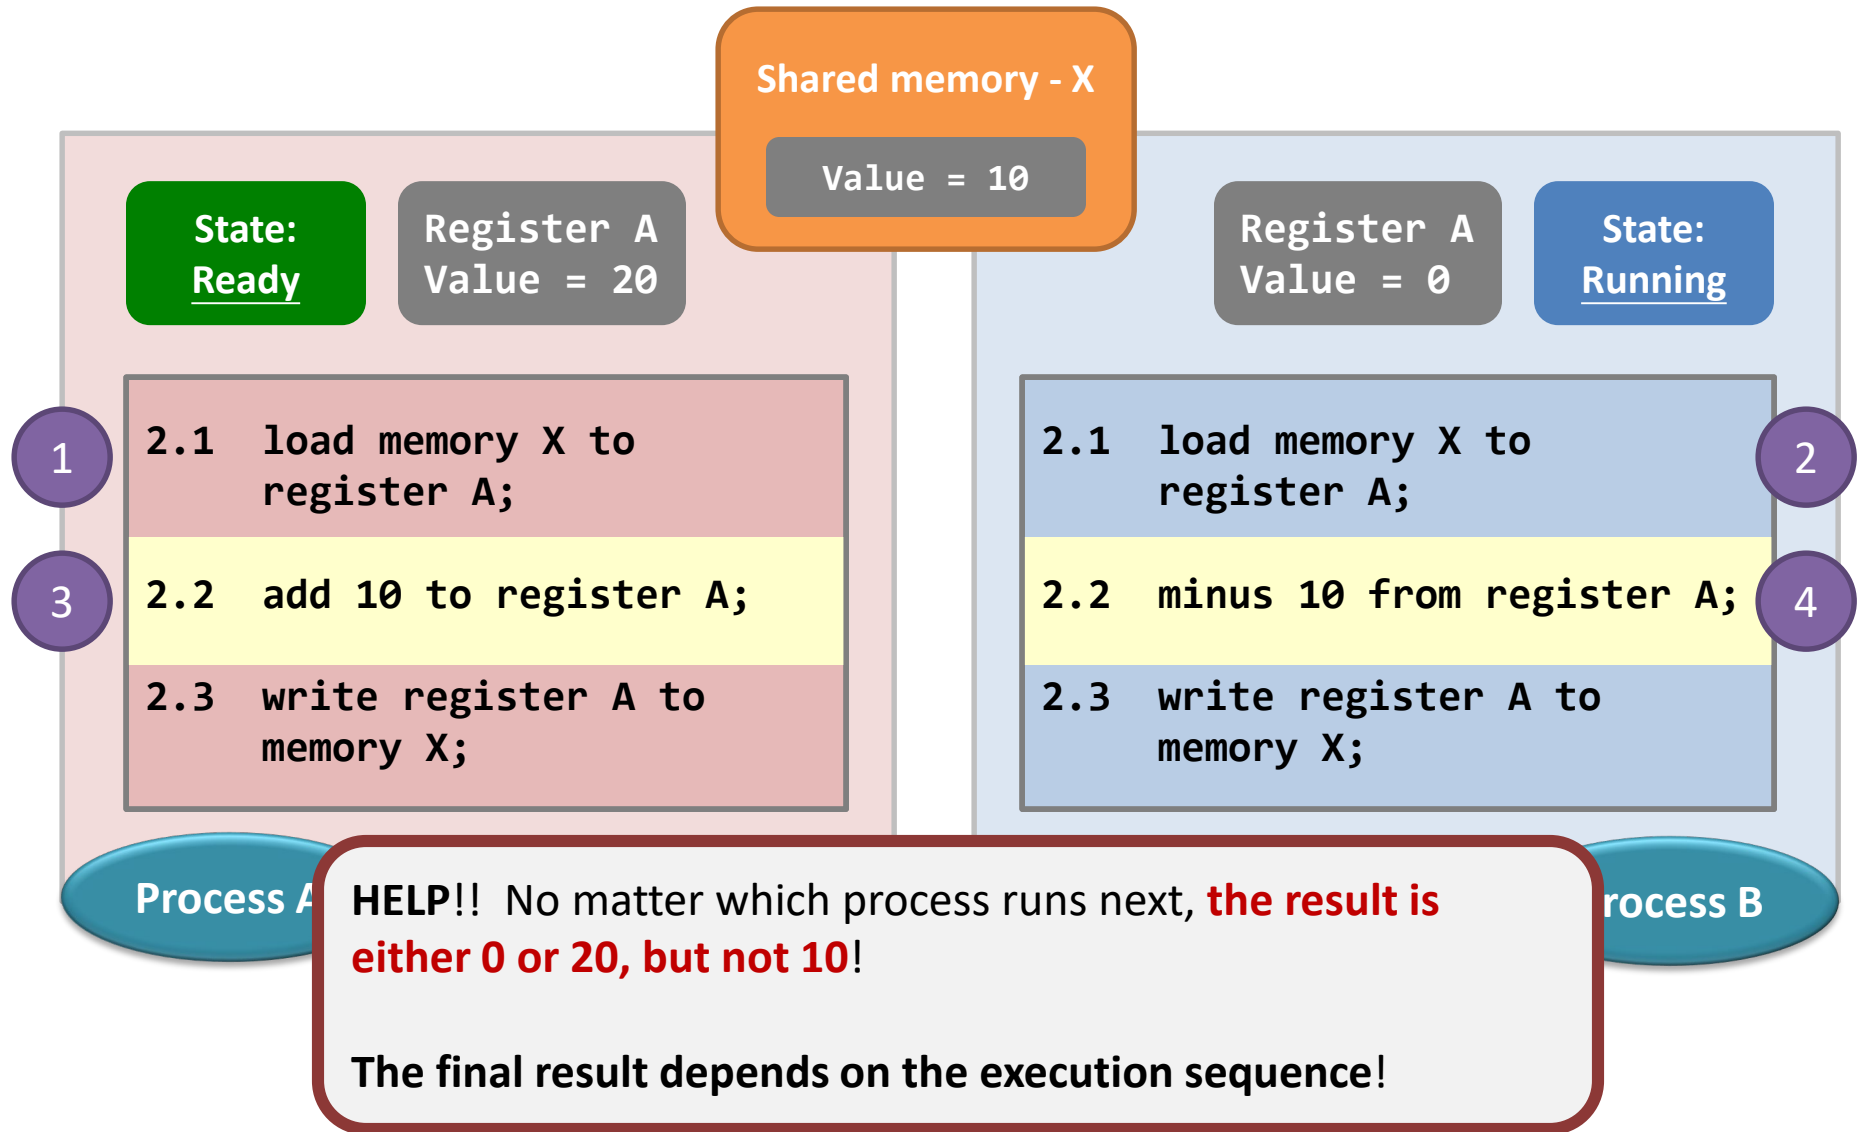
# Problem arise...



# Problem arise...



# Problem arise...

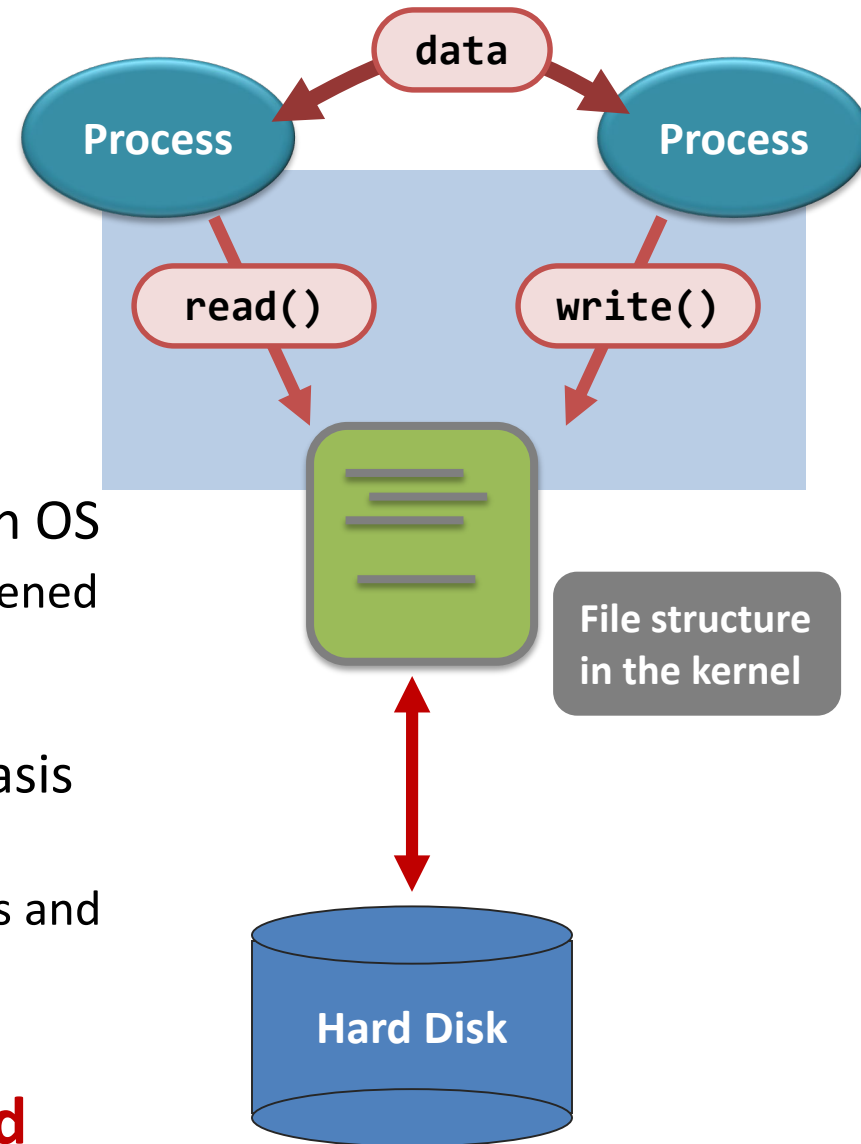


# Race condition – the curse

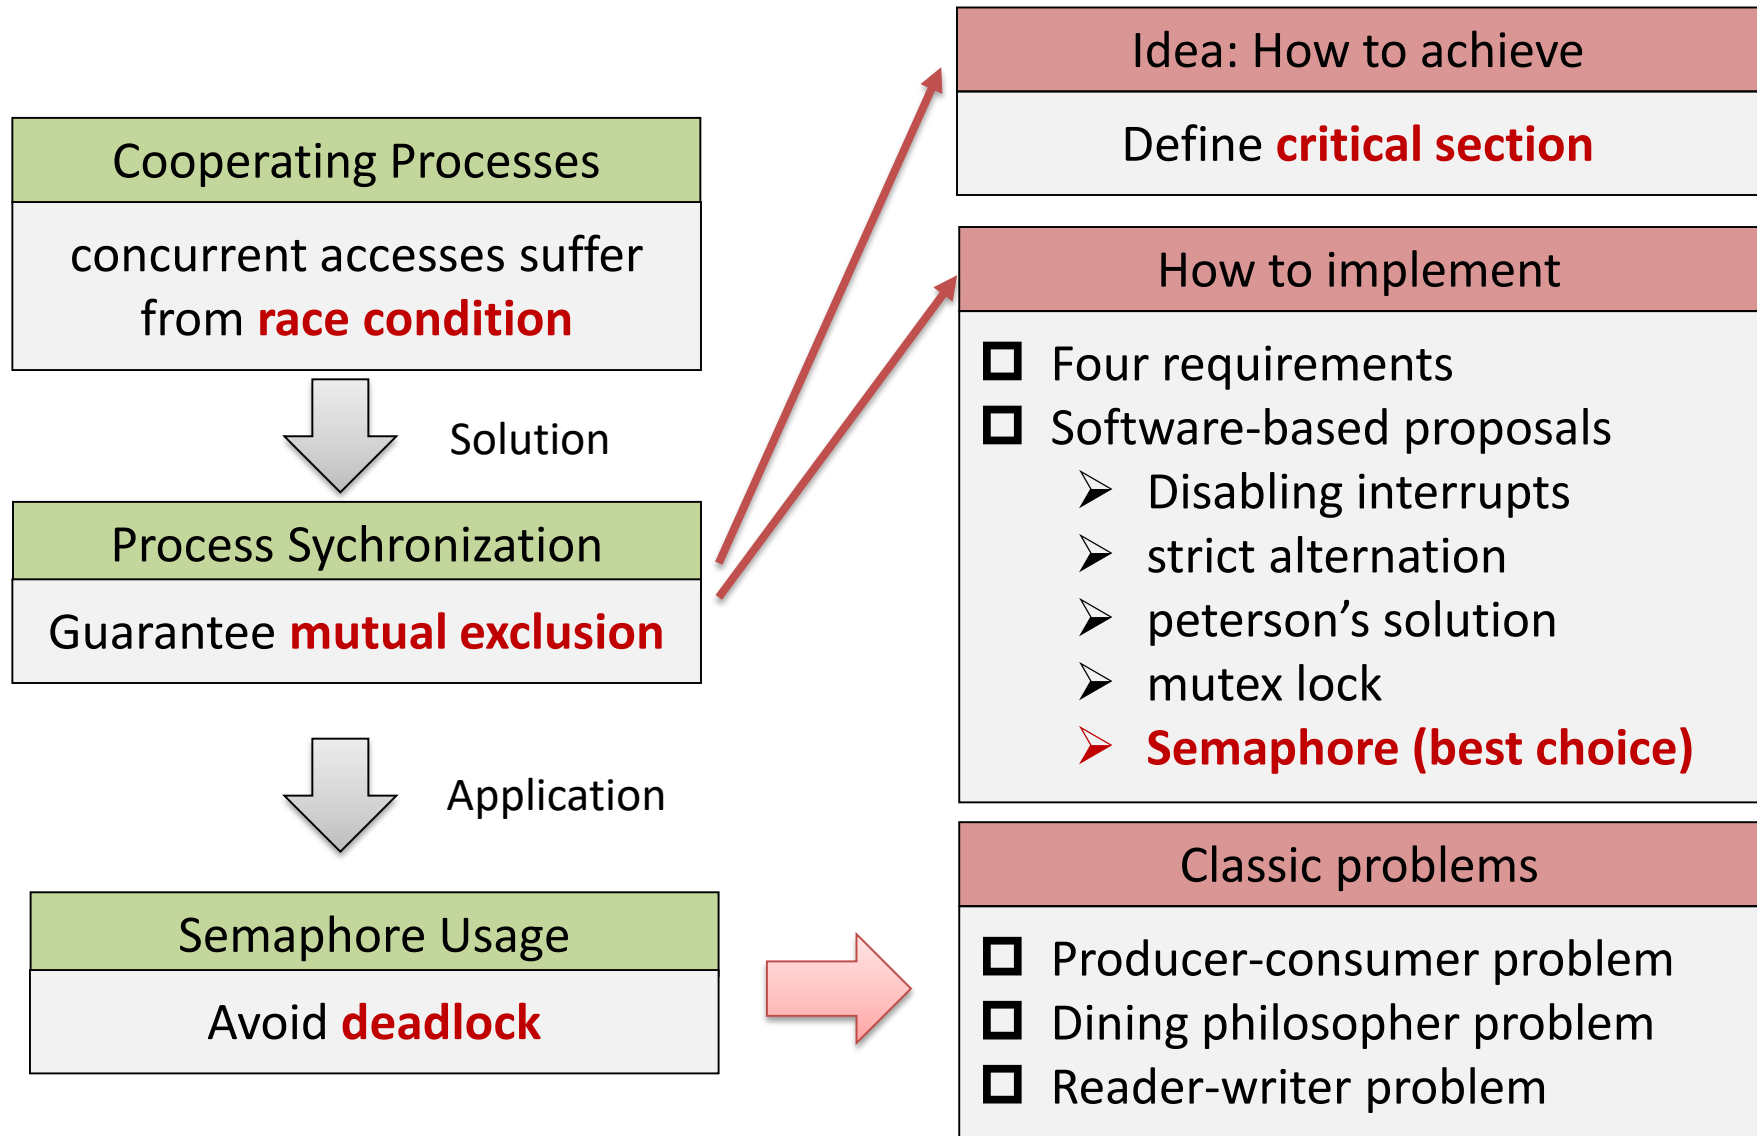
- The above scenario is called the **race condition**.
- A **race condition** means
  - the outcome of an execution depends on a particular order in which the shared resource is accessed.
- Remember: race condition is always a bad thing and debugging race condition has no fun at all!
  - It may end up ...
    - 99% of the executions are fine.
    - 1% of the executions are problematic.

# Race condition – the curse

- For shared memory and files, **concurrent access may yield unpredictable outcomes**
  - **Race condition**
- Common situation
  - Resource sharing occurs frequently in OS
    - EXP: Kernel DS maintaining a list of opened files, maintaining memory allocation, process lists...
  - Multicore brings an increased emphasis on multithreading
    - Multiple threads share global variables and dynamically allocated memory
- **Process synchronization is needed**




# Topics in Process Synchronization






# Inter-process communication (IPC)

- Mutual exclusion
  - what & how to achieve?



# Mutual Exclusion



Two processes playing with the same shared memory is dangerous.


We will face the curse - **race condition**.

The solution can be simple:

**When I'm playing with the shared memory, no one could touch it.**

This is called **mutual exclusion**.

A set of processes would not have the problem of race condition *if mutual exclusion is guaranteed*.

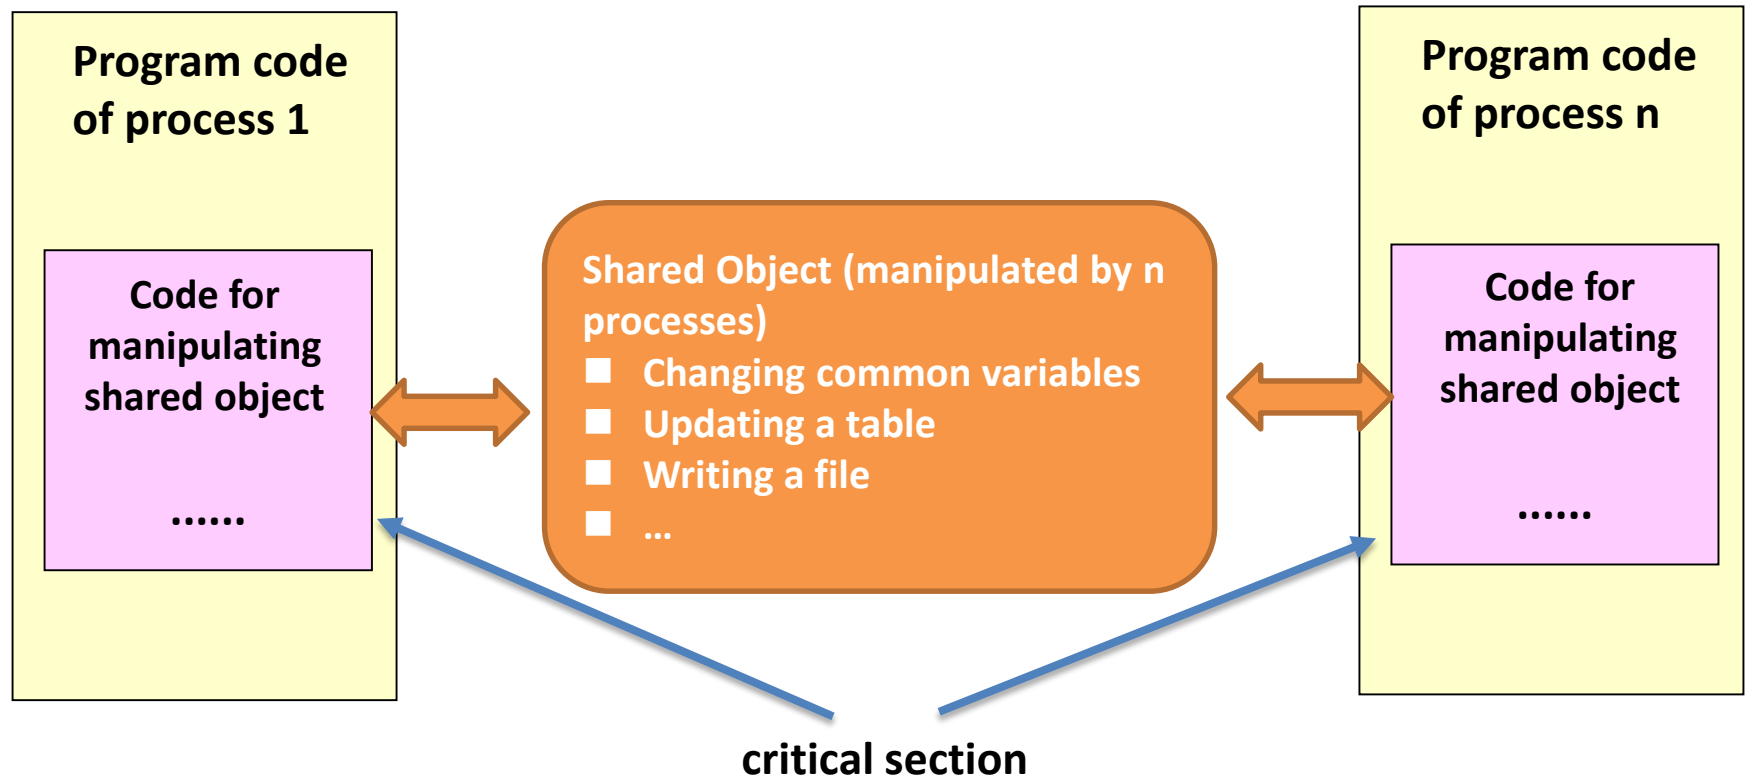


# How to realize mutual exclusion?

- Kernel
  - Preemptive kernels and nonpreemptive kernels
    - Allows (not allow) a process to be preempted while it is running in kernel mode
  - A nonpreemptive kernel is essentially free from race conditions on kernel data structures, and also easy to design (especially for SMP architecture)
  - Why would anyone favor a preemptive kernel
    - More responsive
    - More suitable for real-time programming

# Mutual Exclusion

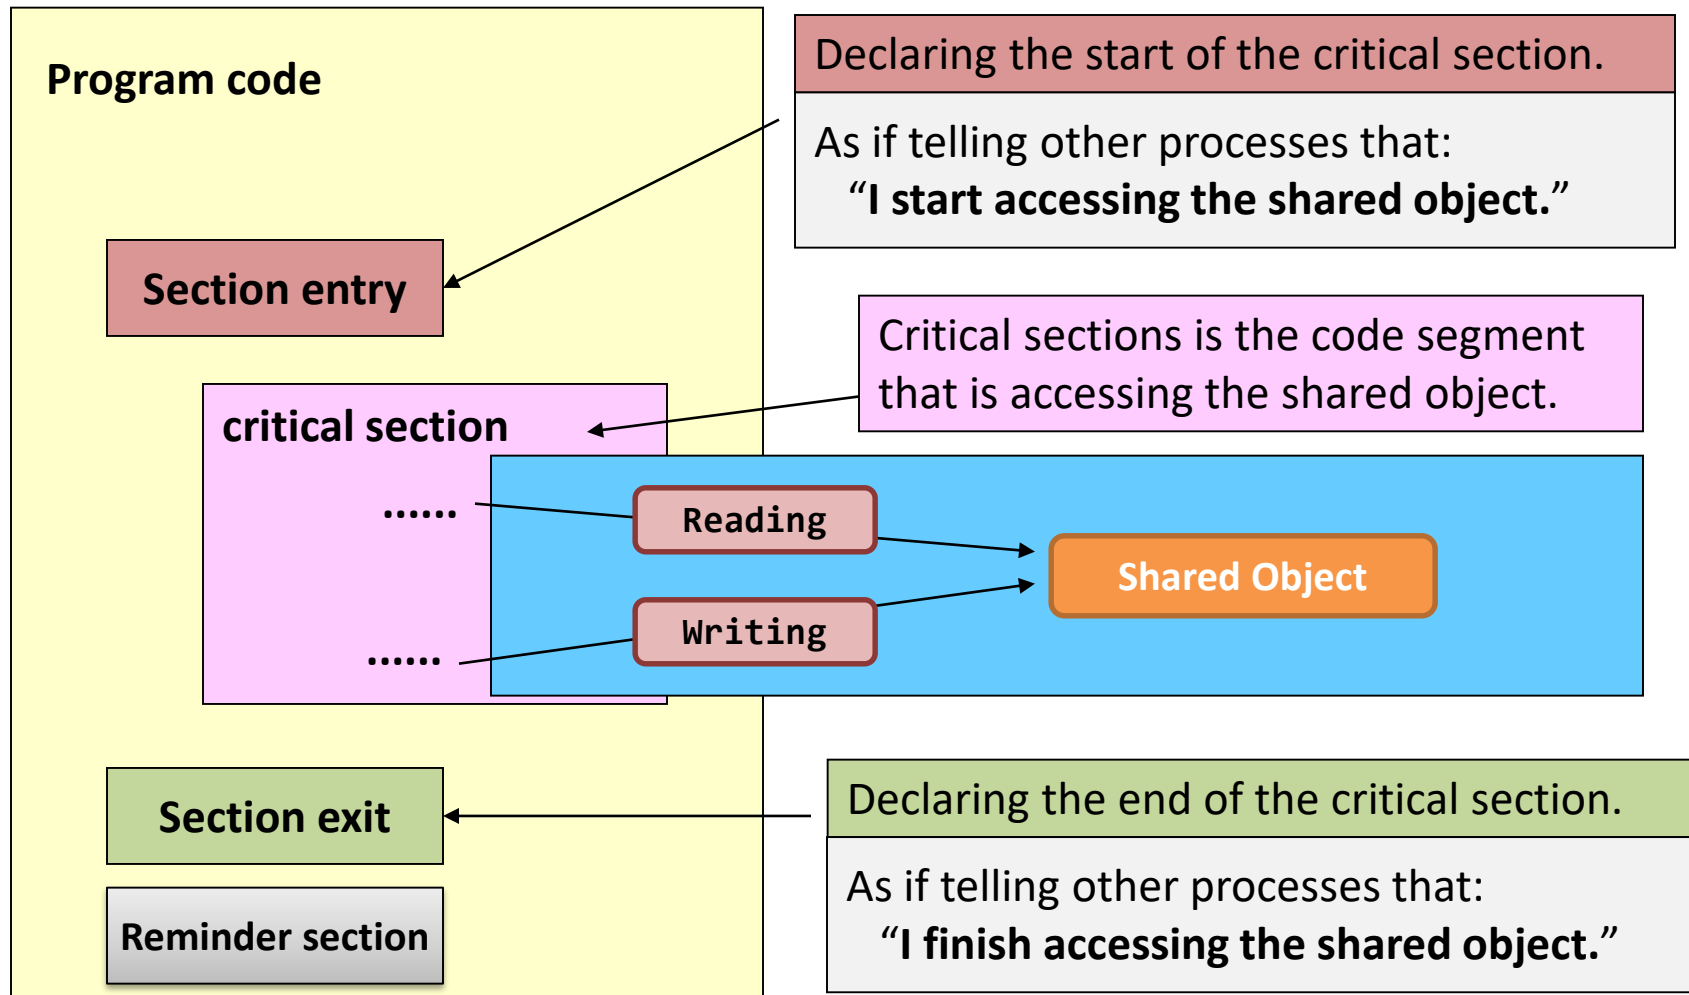
- More generally, how to realize?



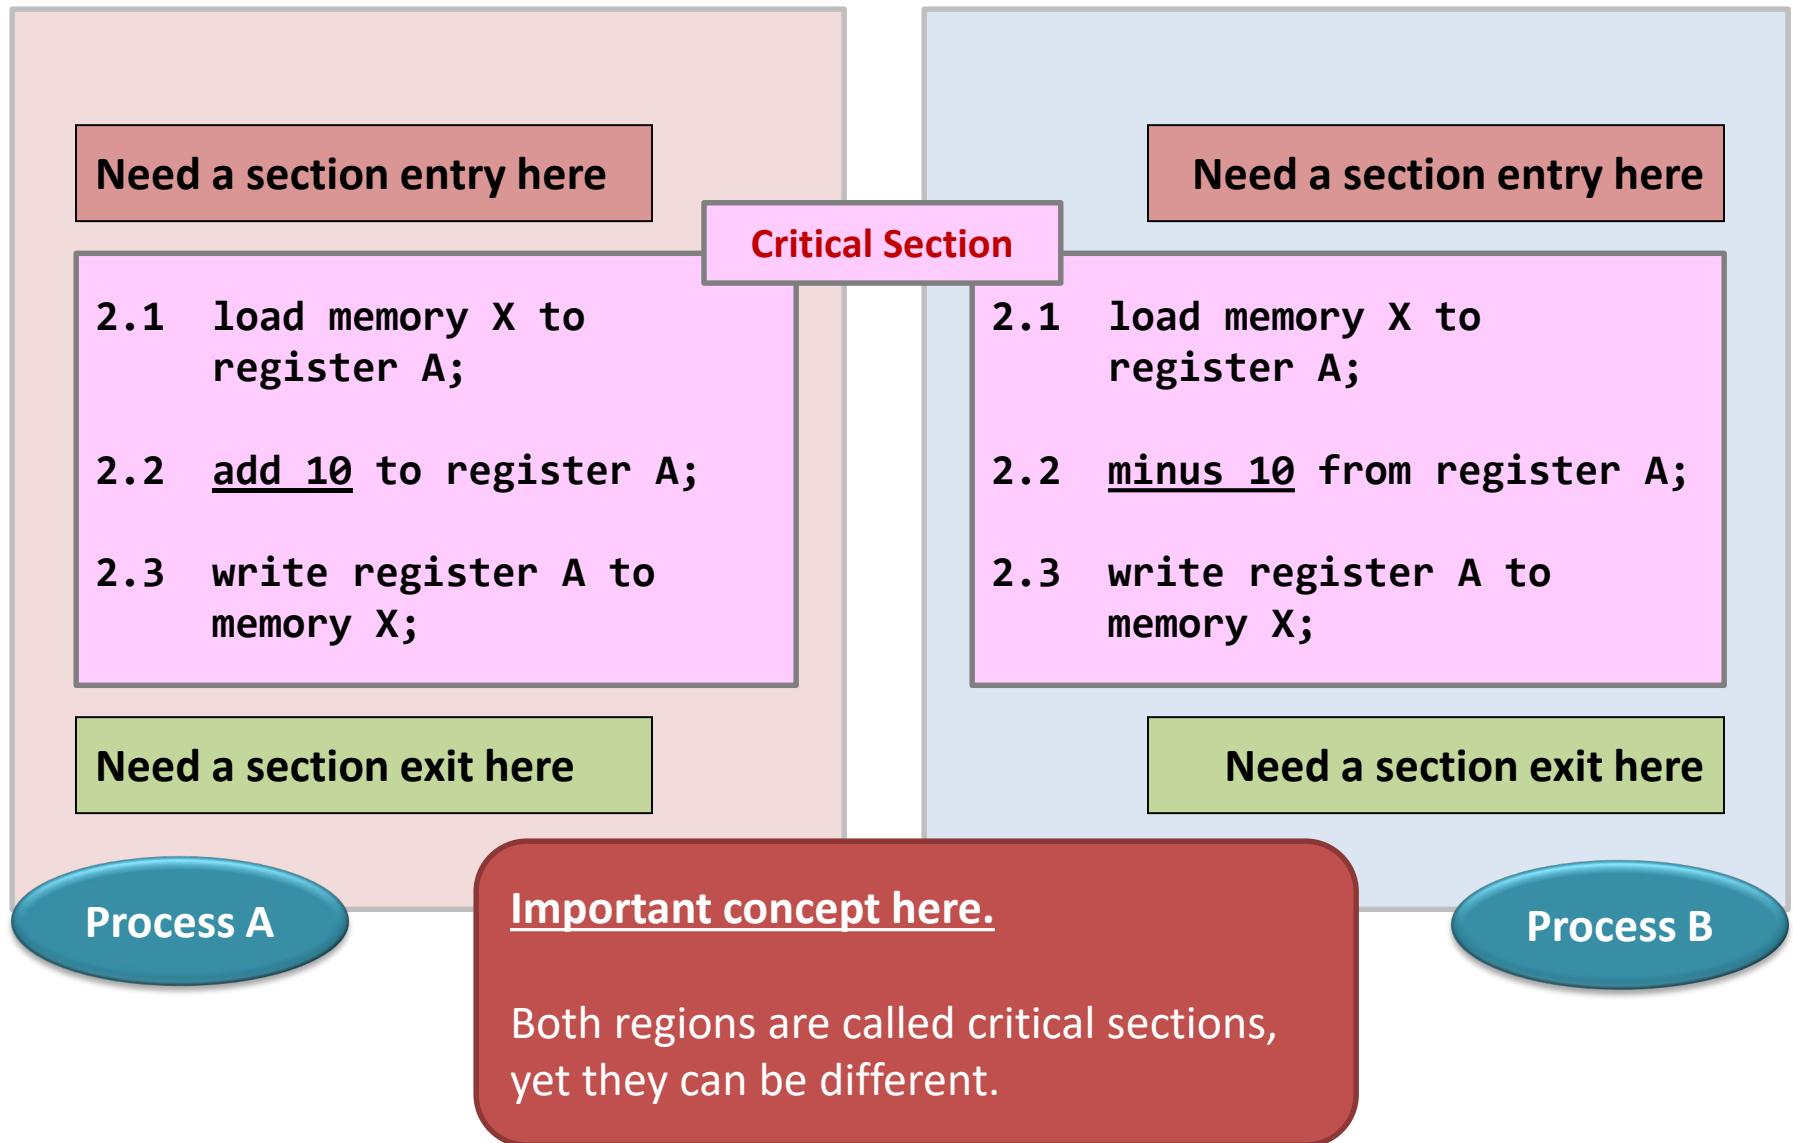
**Solution:** To guarantee that when one process is executing in its critical section, no other process is allowed execute in its critical section.

# Critical Section – General Structure

To guarantee that when one process is executing in its critical section, no other process is allowed execute in its critical section.



# Critical Section – Example



# Summary...for the content so far...

- **Race condition** is a problem.
  - It makes a concurrent program producing **unpredictable** results if you are using shared objects as the communication medium.
  - The outcome of the computation **totally depends on the execution sequences** of the processes involved.
- **Mutual exclusion** is a requirement.
  - If it could be achieved, then the problem of the race condition would be gone.
  - Mutual exclusion hinders the performance of parallel computations.

# Summary...for the content so far...

- **Defining critical sections** is a solution.
  - They are code segments that access shared objects.
  - Critical section must be **as tight as possible**.
    - Well, you can declare the entire code of a program to be a big critical section.
    - But, the program will be a very high chance to block other processes or to be blocked by other processes.
  - Note that one critical section can be designed for **accessing more than one shared objects**.




# Summary...for the content so far...

- **Implementing section entry and exit** is a challenge.
  - The entry and the exit are **the core parts that guarantee mutual exclusion**, but not the critical section.
  - Unless they are correctly implemented, race condition would appear.

# Inter-process communication (IPC)

- Mutual exclusion:
  - how to achieve?
  - how to implement?  
(section entry and exit)



# Entry and exit implementation - requirements

- **Requirement #1: Mutual Exclusion**. No two processes could be simultaneously inside their critical sections.

**Implication**: when one process is inside its critical section, any attempts to go inside the critical sections by other processes are not allowed.

- **Requirement #2**. Each process is executing at a nonzero speed, but no assumptions should be made about the relative speed of the processes and the number of CPUs.

**Implication**: the solution **cannot depend on the time spent inside the critical section**, and the solution cannot assume the number of CPUs in the system.

# Entry and exit implementation - requirements

- **Requirement #3: progress.** No process running outside its critical section should block other processes.

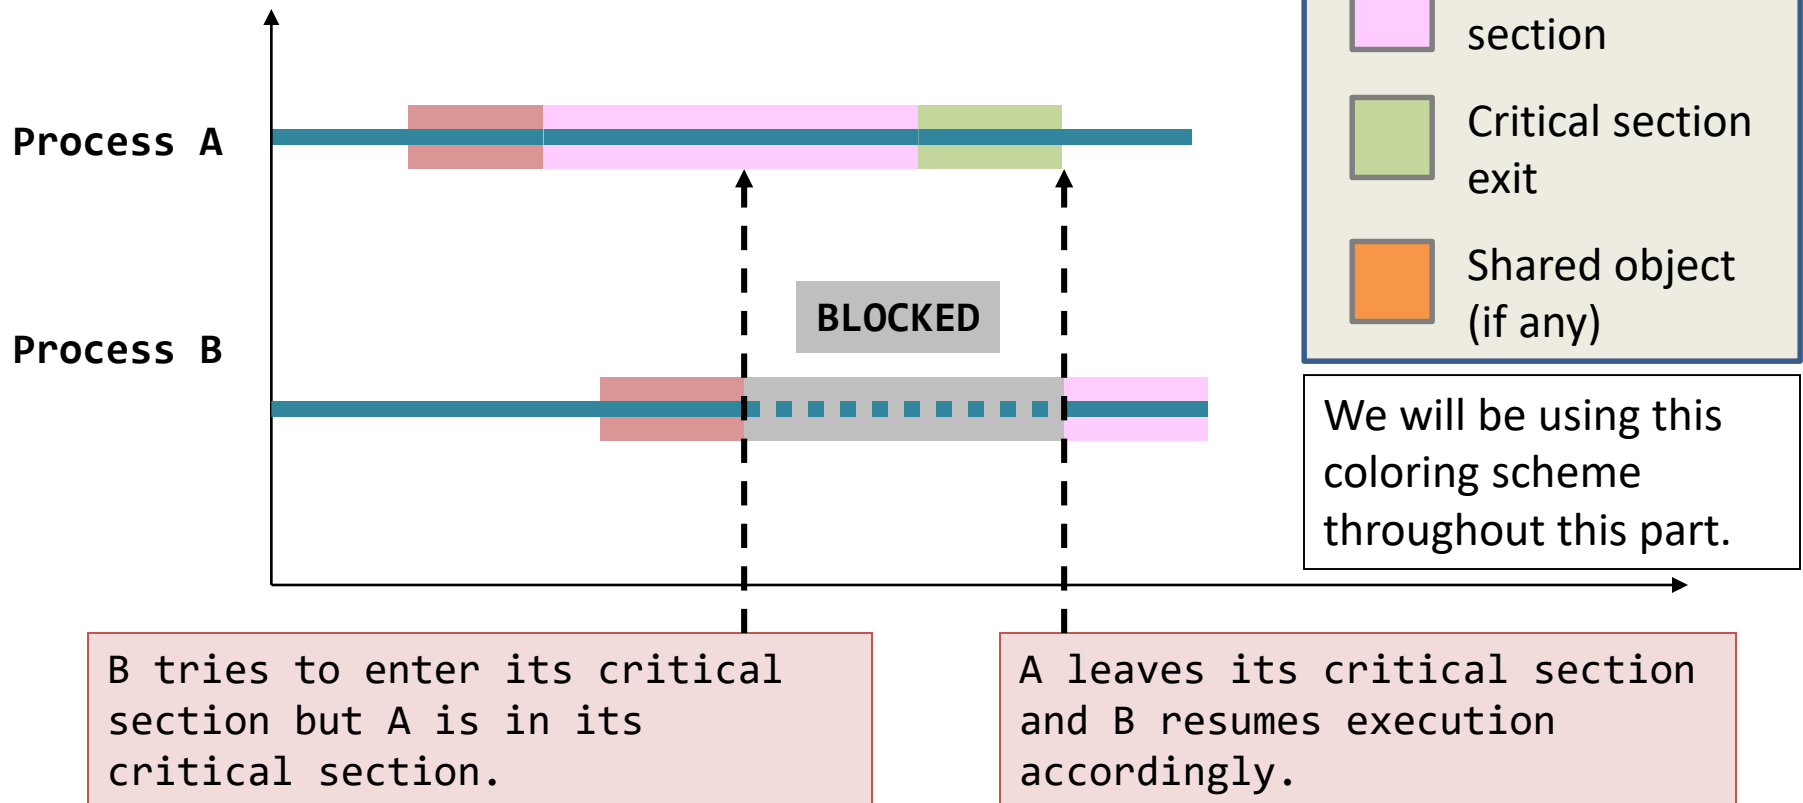
**Implication:** Only processes that are **not executing in their remainder sections** can participate in deciding which will enter its critical section.

- **Requirement #4: Bounded waiting.** No process would have to wait forever in order to enter its critical section.

**Implication:** There exists a bound or limit on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section (no processes should be **starved to death**).

# A typical mutual exclusion scenario

Remember, it is always the entry blocks other processes, but not the critical section.



# Mutual Exclusion Implementation

- Challenges of Implementing **section entry** & **exit**
  - Both operations must be atomic
  - Also need to satisfy the above requirements
  - Performance consideration
- Hardware solution
  - Rely on atomic instructions
  - `test_and_set()`
  - `compare_and_swap`

# Example: test\_and\_set()

- Definition

```
boolean test_and_set(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
  
    return rv;  
}
```

- Mutual exclusion implementation

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
  
    /* remainder section */  
} while (true);
```

# Example: compare\_and\_swap()

- Definition

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

- Mutual exclusion implementation

How to satisfy  
bounded waiting?

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
} while (true);
```



# Enhanced version

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;

    /* critical section */

    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;

    if (j == i)
        lock = false;
    else
        waiting[j] = false;

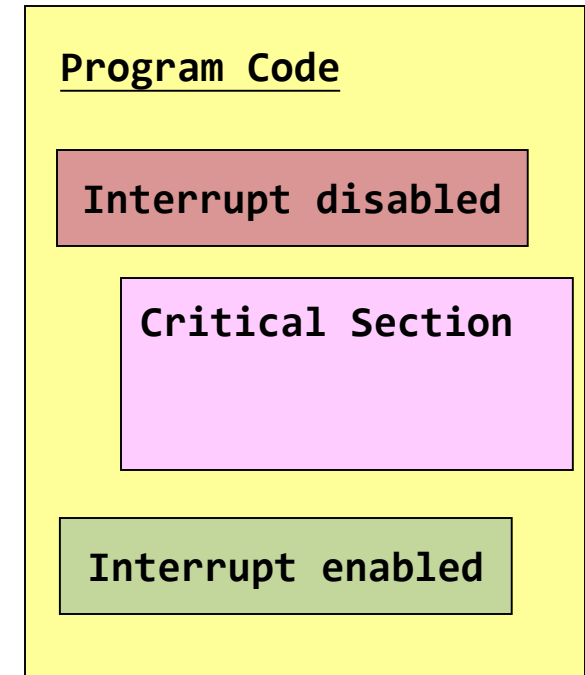
    /* remainder section */
} while (true);
```



lock is initialized as false

# Proposal #1 – disabling interrupt.

- **Method**
  - Similar idea as nonpreemptive kernels
  - To **disable context switching** when the process is inside the critical section.
- **Effect**
  - When a process is in its critical section, no other processes could be able to run.
- **Implementation**
  - A new system call should be provided.
- **Correctness?**
  - **Correct**, but it is not an attractive solution.
  - Not as feasible in a multiprocessor environment
  - Performance issue (may sacrifice concurrency)



# Proposal #2: Mutex Locks

- **Idea**

- A process must acquire the lock before entering a critical section, and release the lock when it exits the critical section
- Using a new shared object to detect the status of other processes, and “**lock**” the shared object

Shared object: “available” (lock)

```
1  acquire(){
2      while(!available)
3          ; /* busy waiting */
4      available = false;
5  }
```

```
1  release(){
2      available = true;
3  }
```

# Proposal #2: Mutex Locks

- **Implementation**

- Calls to acquire and release locks must be performed **atomically**
- Often use hardware instructions

- **Issue**

- Busy waiting: Waste CPU resource
  - **Spinlock**

- **Applications**

- Multiprocessor system
  - When locks are expected to be held for short times

Note that: all processes run the following same code.

Program Code


```
acquire();
```

```
Critical Section
```

```
release();
```

# Other software-based solutions

- Aim
  - To decide which process could go into its critical section

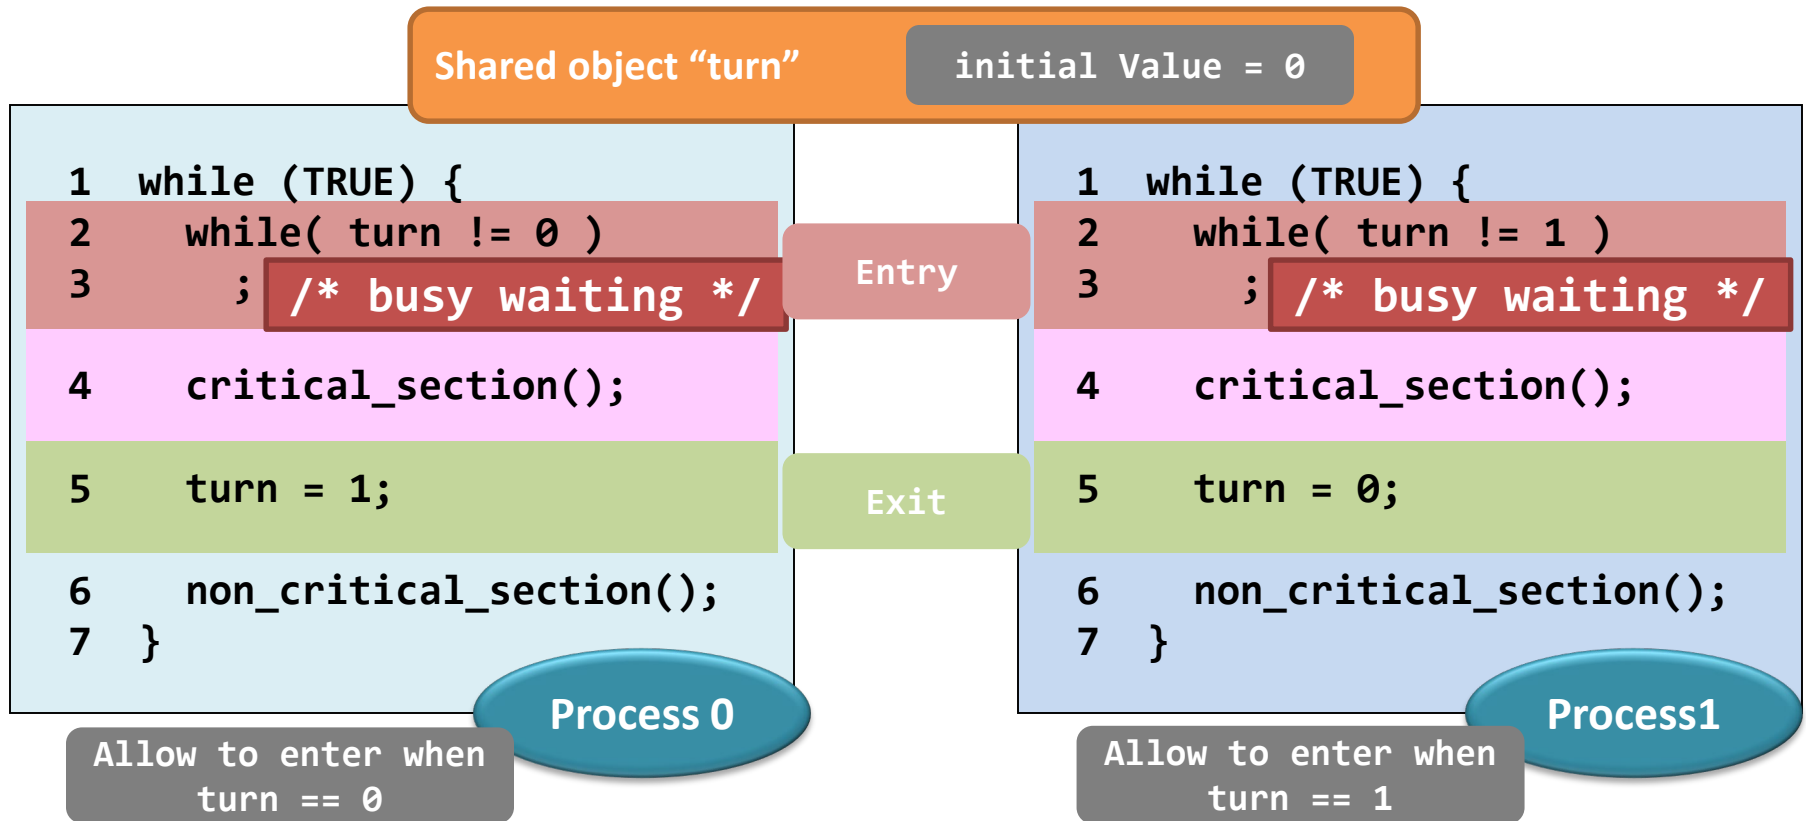


- Key Issues
  - Detect the status of processes (section entry)
    - Need other shared variables
  - Atomicity of section entry and exit


# Proposal #3: Strict alternation

- **Method**

- Using a new shared object to detect the status of other processes



# Proposal #3: Strict alternation



Shared object "turn"

initial Value = 0

```
1 while (TRUE) {
2   while( turn != 0 )
3     ; /* busy waiting */
4   critical_section();
5   turn = 1;
6   non_critical_section();
7 }
```

Process 0

```
1 while (TRUE) {
2   while( turn != 1 )
3     ; /* busy waiting */
4   critical_section();
5   turn = 0;
6   non_critical_section();
7 }
```

Process1

# Proposal #3: Strict alternation - Cons

- Strict alternation seems good, yet, it is **inefficient**.
  - Busy waiting wastes CPU resources.
- In addition, the alternating order is **too strict**.
  - What if Process 0 wants to enter the critical section **twice in a row**? **NO WAY!**
  - Violate any requirement?

**Requirement #3.** No process running outside its critical section should block other processes.



# Proposal #4: Peterson's solution

- How to improve the strict alternation proposal?
  - The Peterson's solution
- Highlights:
  - Share two data items
    - **int turn;** //whose turn to enter its critical section
    - **Boolean interested[2];** //if a process wants to enter
  - Processes would act as a gentleman: if you want to enter, I'll let you first
  - No alternation is there

# Proposal #4: Peterson's solution

Shared object: "turn" &  
"interested[2]"

```
1  int turn;                                /* who can enter critical section */
2  int interested[2] = {FALSE,FALSE};      /* wants to enter critical section*/
3
4  void enter_region( int process ) {       /* process is 0 or 1 */
5      int other;                          /* number of the other process */
6      other = 1-process;                  /* other is 1 or 0 */
7      interested[process] = TRUE;         /* want to enter critical section */
8      turn = other;
9      while ( turn == other &&
              interested[other] == TRUE )
10         ; /* busy waiting */
11 }
12
13 void leave_region( int process ) {       /* process: who is leaving */
14     interested[process] = FALSE;         /* I just left critical region */
15 }
```

Entry

Exit

# Proposal #4: Peterson's solution

```
1  int turn;
2  int interested[2] = {FALSE,FALSE};
3
4  void enter_region( int process ) {
5      int other;
6      other = 1-process;
7      interested[process] = TRUE;
8      turn = other;
9      while ( turn == other &&
              interested[other] == TRUE
10          ;      /* busy waiting */
11  }
12
13 void leave_region( int process ) {
14     interested[process] = FALSE;
15 }
```

Line 8 therefore makes the other one the turn to run.

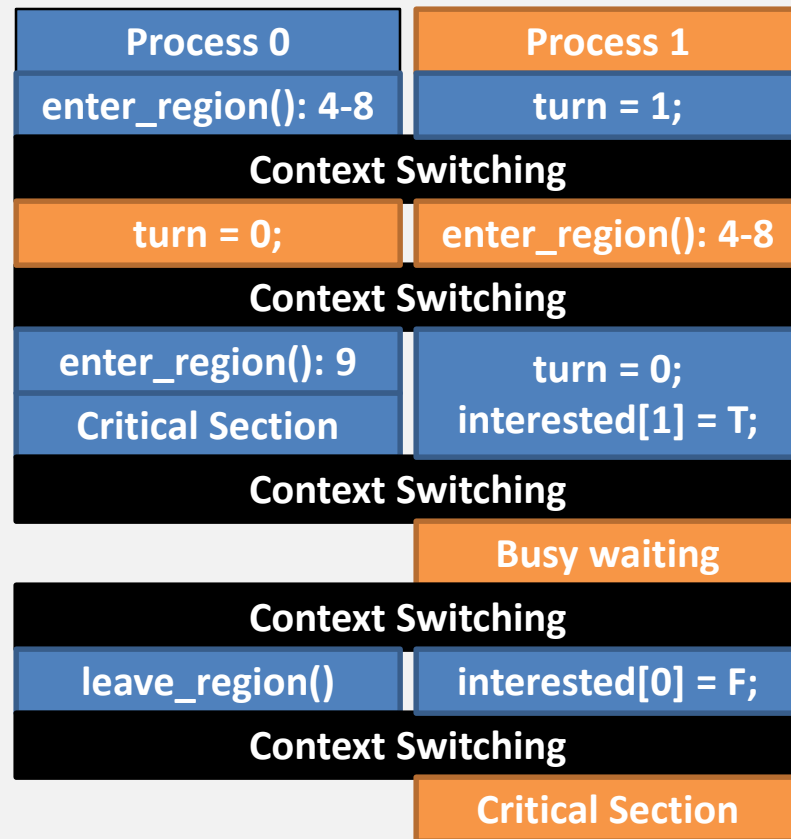
Of course, the process is willing to wait when she wants to enter the critical section.

*"I'm a gentleman!"*

The process always let another process to enter the critical region first although she wants to enter too.

# Proposal #4: Peterson's solution

```
1  int turn;
2  int interested[2] = {FALSE,FALSE};
3
4  void enter_region( int process ) {
5      int other;
6      other = 1-process;
7      interested[process] = TRUE;
8      turn = other;
9      while ( turn == other &&
              interested[other] == TRUE )
10         ;    /* busy waiting */
11 }
12
13 void leave_region( int process ) {
14     interested[process] = FALSE;
15 }
```

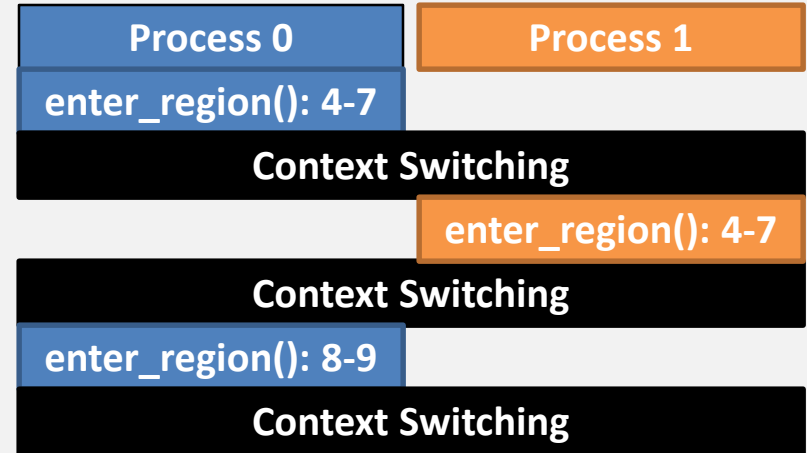


and the story goes on...

Can you show that the requirements are satisfied?

# Proposal #4: Peterson's solution

```
1  int turn;
2  int interested[2] = {FALSE,FALSE};
3
4  void enter_region( int process ) {
5      int other;
6      other = 1-process;
7      interested[process] = TRUE;
8      turn = other;
9      while ( turn == other &&
10             interested[other] == TRUE )
11          ;    /* busy waiting */
12
13 void leave_region( int process ) {
14     interested[process] = FALSE;
15 }
```

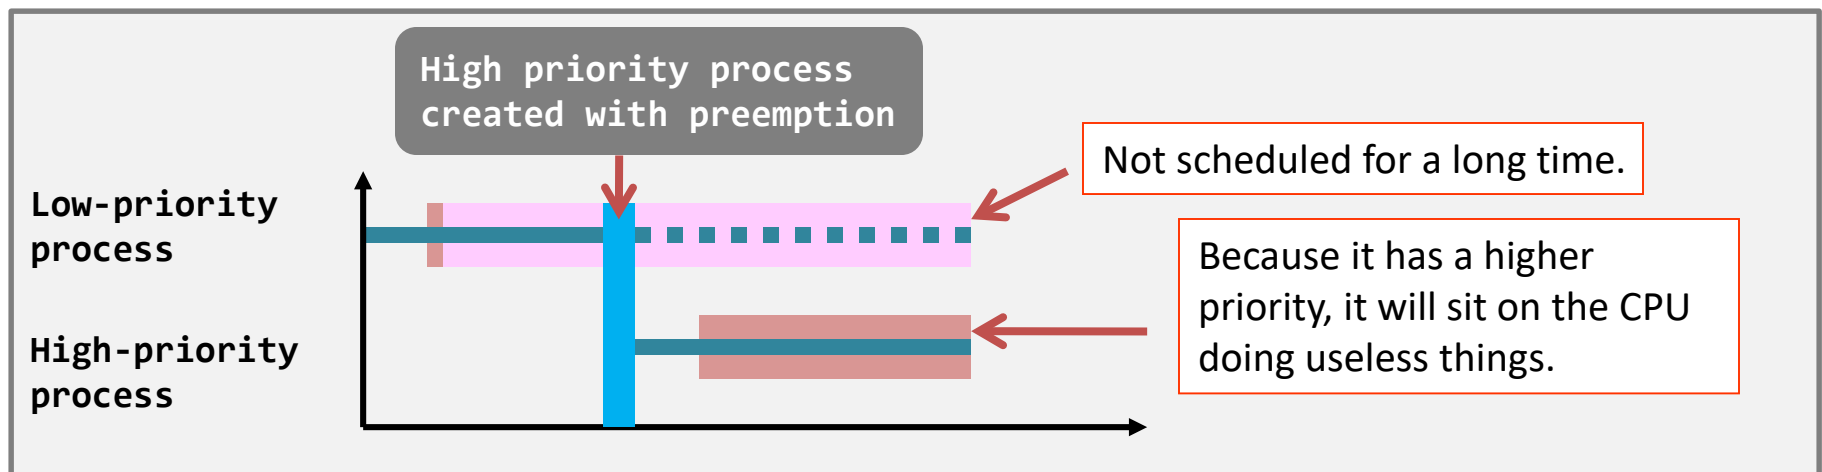


Can you complete the flow?  
(what is the difference?)

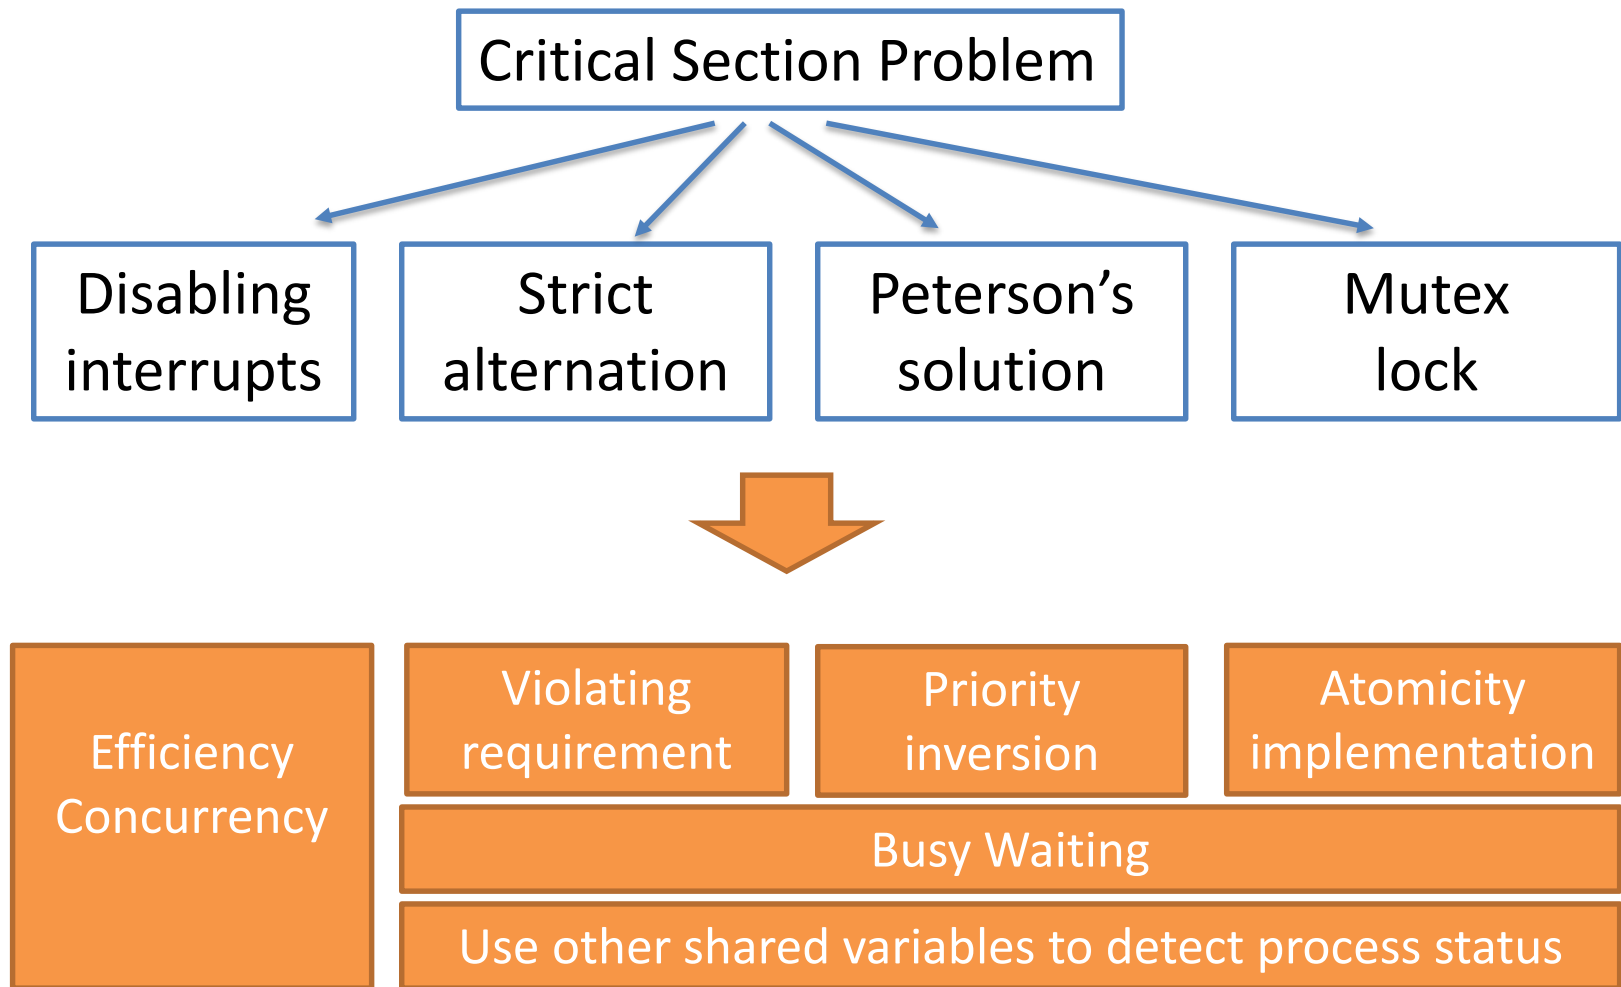
Can both processes progress?

# Proposal #4: Peterson's solution – issues

- Busy waiting has its own problem...
  - An apparent problem: wasting CPU time.
  - A hidden, serious problem: **priority inversion problem**.
    - A low priority process is inside the critical region, but ...
    - A high priority process wants to enter the critical region.
    - Then, the high priority process will perform busy waiting for a long time or even forever.




# Story so far...



# Final proposal: Semaphore

- In real life, semaphore is a flag signaling system.
  - It tells a train driver (or a plane pilot) when to stop and when to proceed.
- When it comes to programming...
  - A semaphore is a data type.
  - You can imagine that it is **an integer** (but it is certainly not an integer when it comes to real implementation).



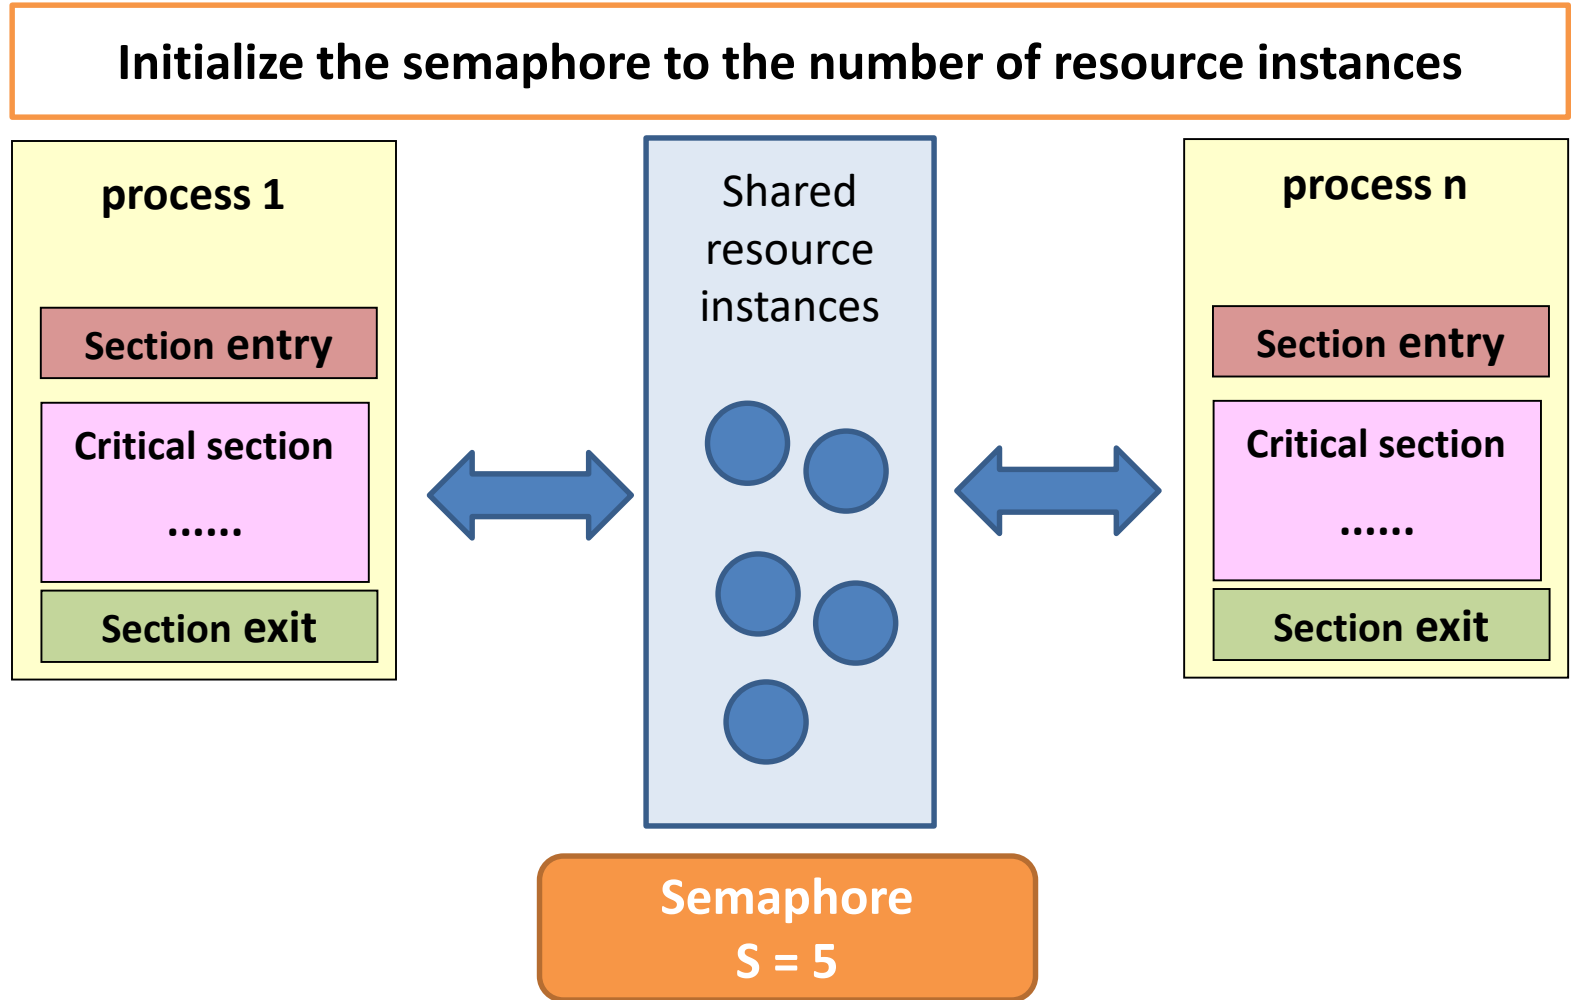


# Final proposal: Semaphore

- Semaphore is a data type (**additional shared object**)
  - Accessed only through two standard **atomic** operations
  - **down()**: originally termed P (from Dutch *proberen*, “to test”), **wait()** in textbook
    - Decrementing the count
  - **up()**: originally termed V (from *verhogen*, “to increment”), **signal()** in textbook
    - Incrementing the count
- Two types
  - **Binary semaphore**: 0 or 1 (similar to mutex lock)
  - **Counting semaphore**: control finite number of resources

# Final proposal: Semaphore


- Idea



# Final proposal: Semaphore

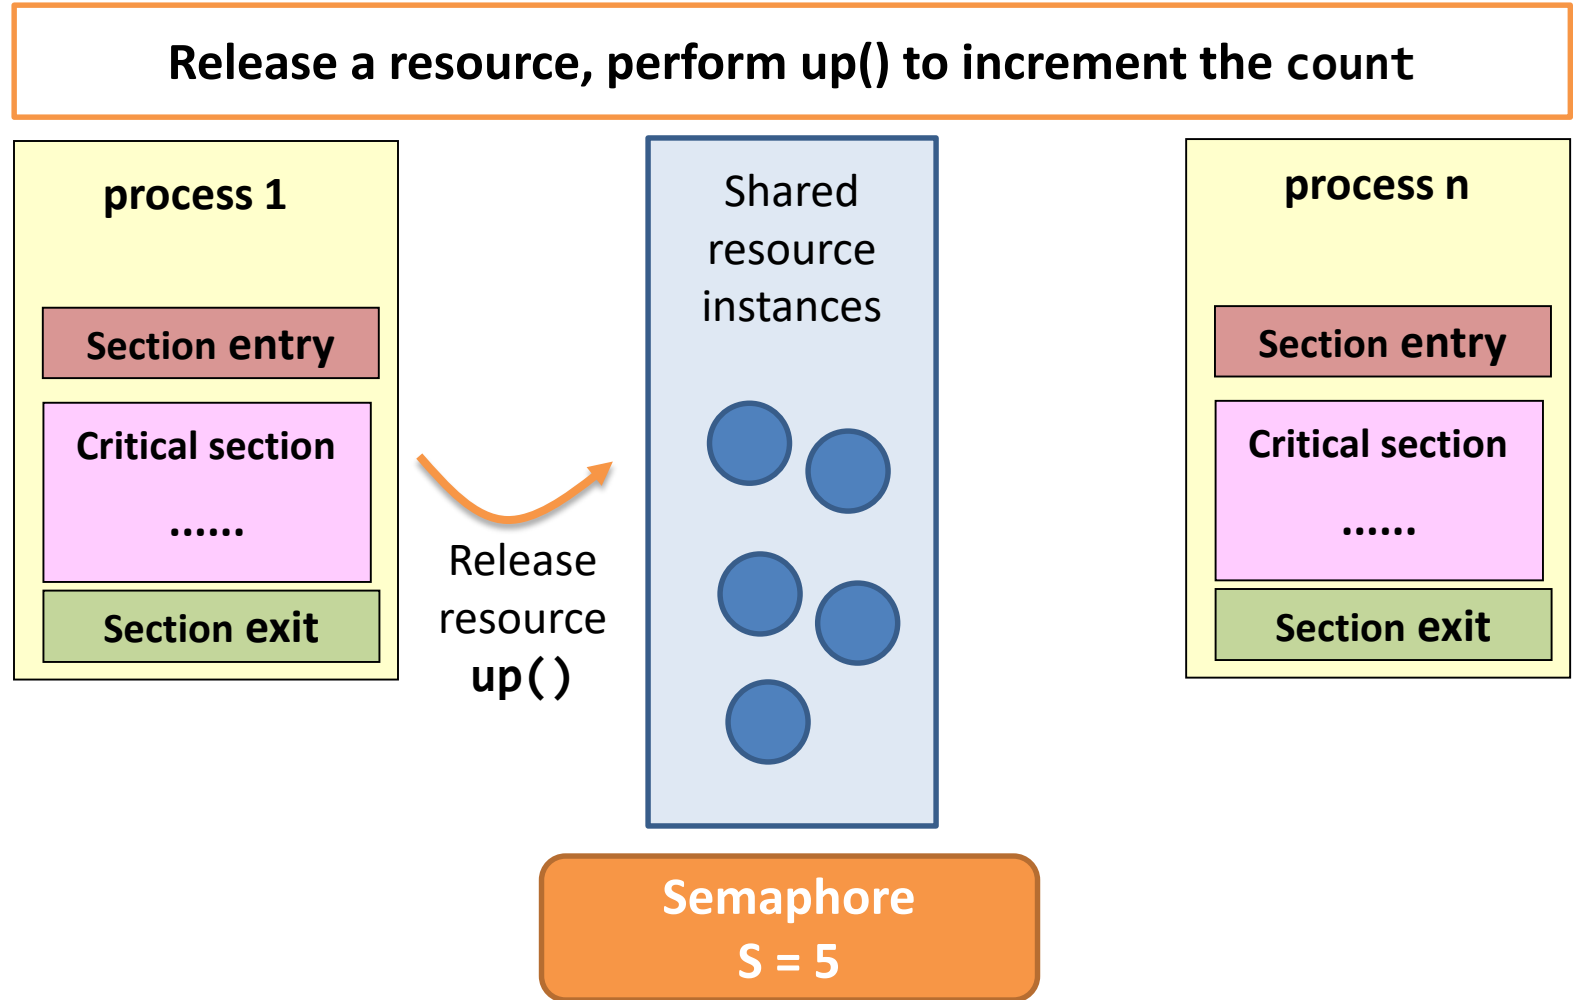
- Idea

Wish to use a resource, perform `down()` to decrement the count



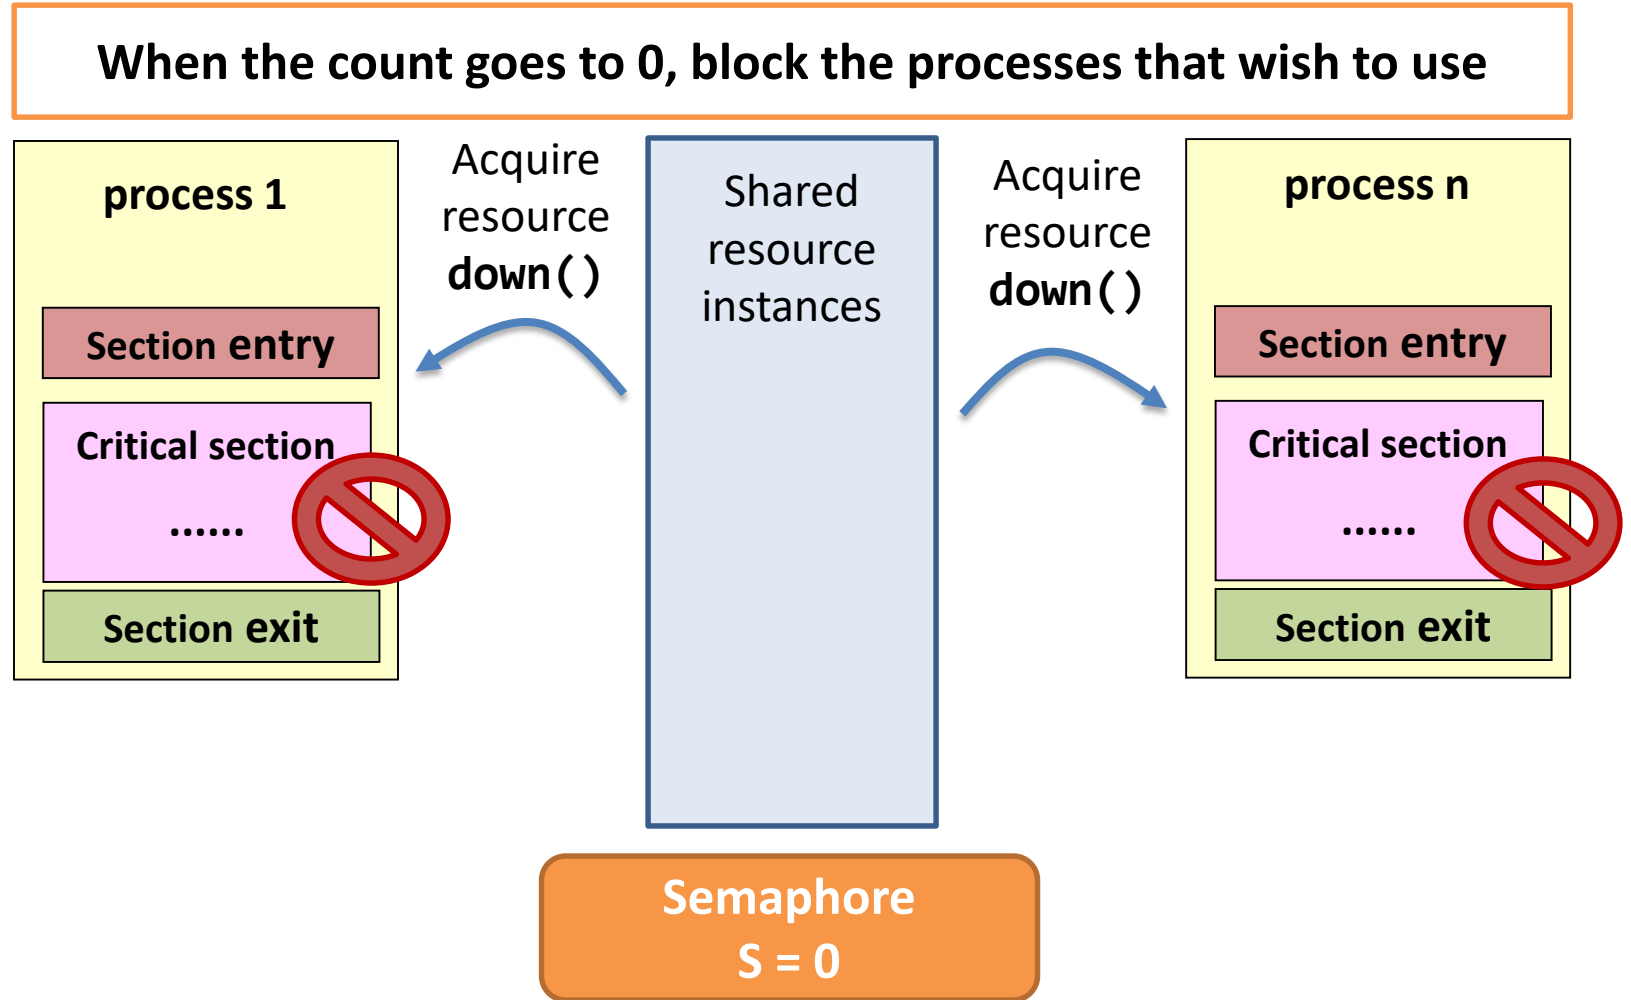
# Final proposal: Semaphore

- Idea



# Final proposal: Semaphore

- Idea



# Semaphore – Simple Implementation

## Data Type definition

```
typedef int semaphore;
```

**Counting Semaphore:** initialized to be the number of resources available

## Section Entry: down()

```
1 void down(semaphore *s) {  
2  
3     while ( *s == 0 ) {  
4  
5         ;//busy wait  
6  
7     }  
8     *s = *s - 1;  
9  
10 }
```

## Section Exit: up()

```
1 void up(semaphore *s) {  
2  
3  
4  
5     *s = *s + 1;  
6  
7 }
```

# Semaphore – Address busy waiting

## Data Type definition

```
typedef int semaphore;
```

## First issue: Busy waiting

**Solution:** block the process instead of busy waiting (place the process into a waiting queue)

## Section Entry: down()

```
1 void down(semaphore *s) {  
2  
3     while ( *s == 0 ) {  
4  
5         special_sleep();  
6  
7     }  
8     *s = *s - 1;  
9  
10 }
```

## Section Exit: up()

```
1 void up(semaphore *s) {  
2  
3     if ( *s == 0 )  
4         special_wakeup();  
5     *s = *s + 1;  
6  
7 }
```

# Semaphore – Address busy waiting

## Data Type definition

```
typedef int semaphore;
```

## First issue: Busy waiting

**Solution:** block the process instead of busy waiting (place the process into a waiting queue)

```
typedef struct{  
  
    int value;  
    struct process * list;  
}semaphore;
```

## Note

**Implementation:** The waiting queue may be associated with the semaphore, so a semaphore is not just an integer



# Semaphore – Atomicity

## Data Type definition

```
typedef int semaphore;
```

## Section Entry: down()

```
1 void down(semaphore *s) {  
2  
3     while ( *s == 0 ) {  
4  
5         special_sleep();  
6  
7     }  
8     *s = *s - 1;  
9  
10 }
```

**Second issue: Atomicity** (both operations must be atomic)

**Solution:** Disabling interrupts

## Section Exit: up()

```
1 void up(semaphore *s) {  
2  
3     if ( *s == 0 )  
4         special_wakeup();  
5     *s = *s + 1;  
6  
7 }
```

# Semaphore – Atomicity

## Data Type definition

```
typedef int semaphore;
```

## Section Entry: down()

```
1 void down(semaphore *s) {  
2     disable_interrupt();  
3     while ( *s == 0 ) {  
4         enable_interrupt();  
5         special_sleep();  
6         disable_interrupt();  
7     }  
8     *s = *s - 1;  
9     enable_interrupt();  
10 }
```

**Second issue: Atomicity** (both operations must be atomic)

**Solution:** Disabling interrupts

Also, only one process can invoke “**disable\_interrupt()**”. Later processes would be blocked until “**enable\_interrupt()**” is called.

## Section Exit: up()

```
1 void up(semaphore *s) {  
2     disable_interrupt();  
3     if ( *s == 0 )  
4         special_wakeup();  
5     *s = *s + 1;  
6     enable_interrupt();  
7 }
```

# Semaphore – The code

## Data Type definition

```
typedef int semaphore;
```

## Section Entry: down()

```
1 void down(semaphore *s) {  
2     disable_interrupt();  
3     while ( *s == 0 ) {  
4         enable_interrupt();  
5         special_sleep();  
6         disable_interrupt();  
7     }  
8     *s = *s - 1;  
9     enable_interrupt();  
10 }
```

Why need these two statements?

Disabling interrupts may sacrifice concurrency, so it is essential to keep the critical section as short as possible

## Section Exit: up()

```
1 void up(semaphore *s) {  
2     disable_interrupt();  
3     if ( *s == 0 )  
4         special_wakeup();  
5     *s = *s + 1;  
6     enable_interrupt();  
7 }
```

# Semaphore – details

Process 1234

down(X)

Section Entry: down()

```
1 void down(semaphore *s) {  
2     disable_interrupt();  
3     while ( *s == 0 ) {  
4         enable_interrupt();  
5         special_sleep();  
6         disable_interrupt();  
7     }  
8     *s = *s - 1;  
9     enable_interrupt();  
10 }
```

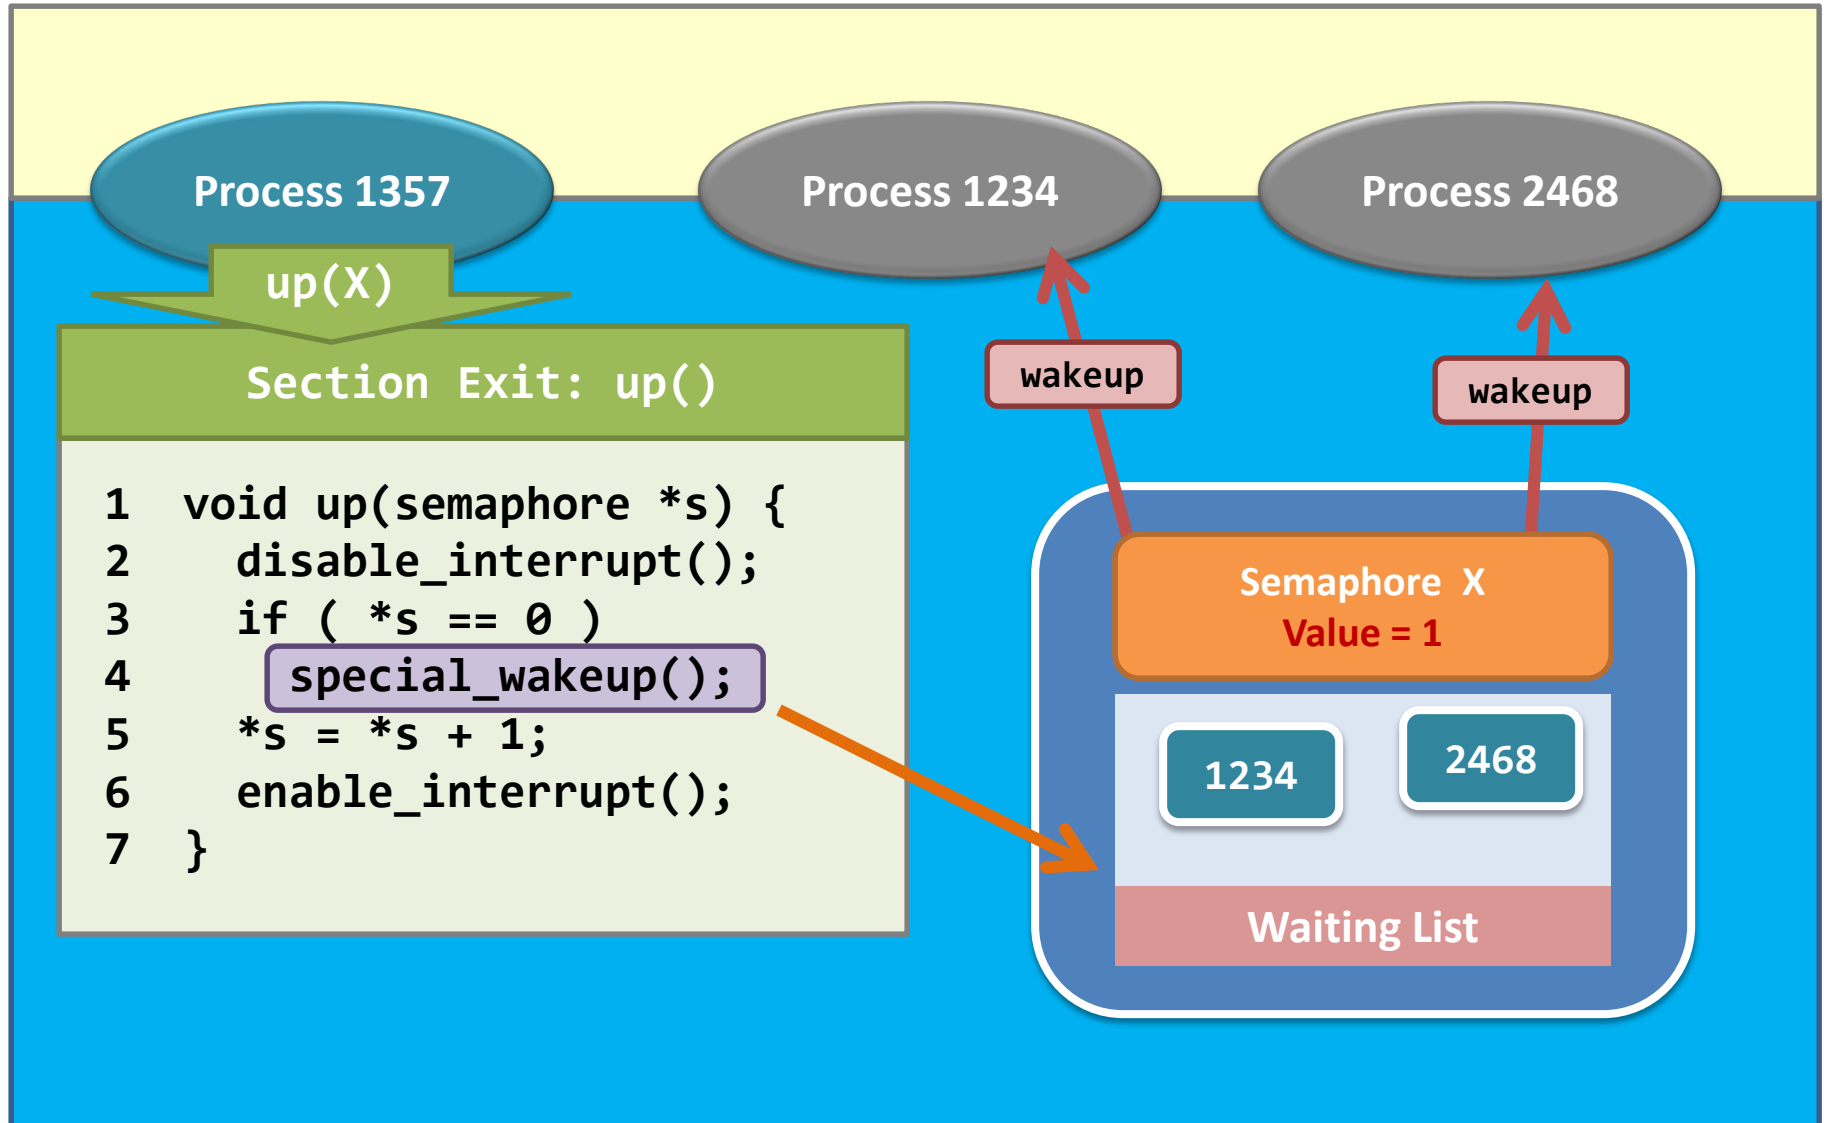
Suppose that process 1234 is willing to access the shared resource (enter its critical section), but no resource is available

Semaphore X  
Value = 0

1234

Waiting List

# Semaphore – details



# Semaphore – details

Process 1234

down(X)

Process 2468

down(X)

Note that it is impossible for **two blocked processes to get out of the down() simultaneously.**

Why?

Only one process can invoke **disable\_interrupt()**

Only one process can manipulate this shared variable

## Section Entry: down()

```
1 void down(semaphore *s) {  
2     disable_interrupt();  
3     while ( *s == 0 ) {  
4         enable_interrupt();  
5         special_sleep();  
6         disable_interrupt();  
7     }  
8     *s = *s - 1;  
9     enable_interrupt();  
10 }
```

here

# Semaphore – details

Process 1234

down(X)

Process 2468

down(X)

Note that it is impossible for **two processes to get out of the down() simultaneously.**

Why?

Whether which process can get out of **down()** is **the business of the scheduler.**

Section Entry: down()


```
1 void down(semaphore *s) {  
2     disable_interrupt();  
3     while ( *s == 0 ) {  
4         enable_interrupt();  
5         special_sleep();  
6         disable_interrupt();  
7     }  
8     *s = *s - 1;  
9     enable_interrupt();  
10 }
```

here

# Semaphore – in action

- Add them together...

Either one of the processes can enter the critical section when the first process calls “up(s)”.



```
semaphore *s;  
*s = 1;      /* initial value */
```

```
1 while(TRUE) {  
2     down(s);  
3     critical_section();  
4     up(s);  
5 }
```

entry

exit



# Summary...on semaphore

- More on semaphore...it demonstrates an important kind of operations – **atomic operations**.

## Definition of atomic operation

- Either none of the instructions of an atomic operation were completed, or
- All instructions of an atomic operation are completed.

- In other words, the entire **up()** and **down()** are indivisible.
  - If it returns, the change must have been made;
  - If it is aborted, no change would be made.

# Summary...on critical section problem

- What happened is just the implementation of mutual exclusion (section entry and section exit).

	Comments
Disabling interrupts	Time consuming for multiprocessor systems, sacrifices concurrency.
Strict alternation	Not a good one, busy waiting & violating one mutual exclusion requirement.
Peterson's solution	Busy waiting & has a potential " <i>priority inversion problem</i> ".
Mutex lock	Busy waiting, often relies on hardware instructions.
Semaphore	<b>BEST CHOICE.</b>

# Story so far...

- Cooperating processes
  - IPC mechanisms (shared memory, pipes, FIFOs, sockets)
  - Race condition
- Synchronization
  - Mutual exclusion
    - Critical section problem
    - Disabling interrupts, strict alternation, Peterson's solution, mutex lock, semaphore
- What is next?
  - How to use semaphore to solve classic IPC problems
  - Deadlock