

# Operating Systems

Associate Prof. Yongkun Li

中科大-计算机学院 副教授

<http://staff.ustc.edu.cn/~ykli>

Ch5

## Process Communication & Synchronization

# Story so far...

- Process concept + operations
  - Programmer's perspective + kernel's perspective
- Thread
  - Lightweight process
- We mainly talked about the stuffs related to a single process/thread, what if multiple processes exist...

# Processes

- The processes within a system may be
  - ***independent*** or
    - Independent process cannot affect or be affected by other processes
  - ***cooperating***
    - Cooperating process can affect or be affected by other processes
- Note: Any process that shares data with others is a cooperating process

# Cooperating Processes


- Why we need cooperating processes
  - Information sharing
    - e.g., shared file
  - Computation speedup
    - executing subtasks in parallel
  - Modularity
    - dividing system functions into separate processes
  - Convenience

# Cooperating Processes

- Paradigm for cooperating processes
  - **Producer-consumer problem**, useful metaphor for many applications (abstracted problem model)
    - *producer* process produces information that is consumed by a *consumer* process
    - At least one producer and one consumer
- Cooperating processes need
  - **interprocess communication (IPC)** for exchanging data

# Inter-process communication (IPC)

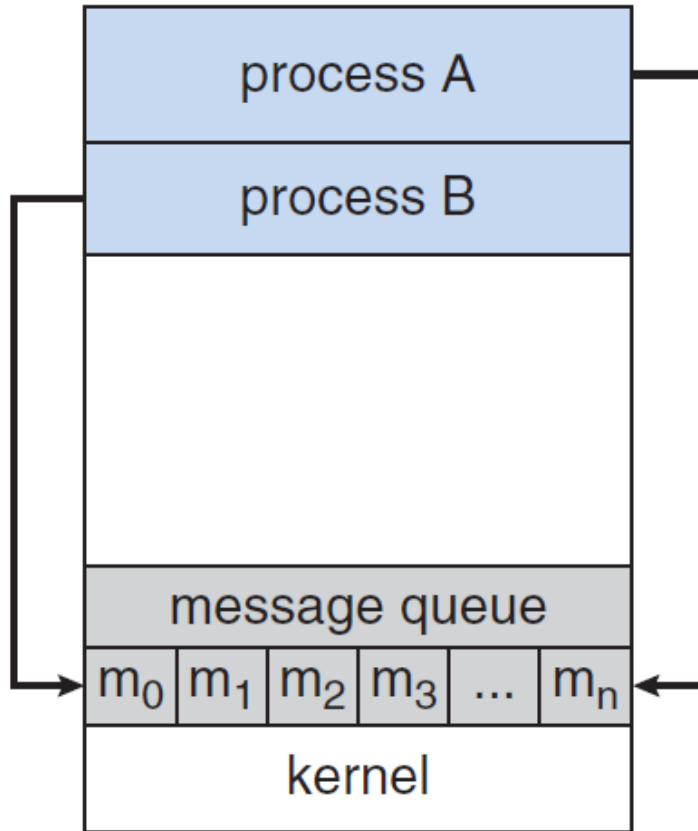
- What and how?



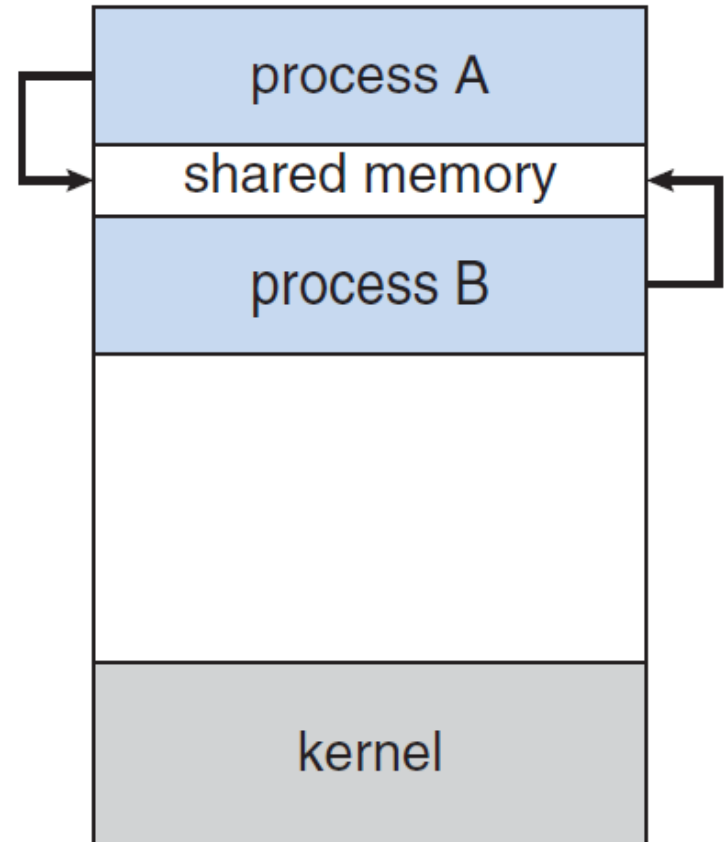
# Interprocess Communication

- IPC: used for exchanging data between processes
- Two (abstracted) models of IPC
  - **Shared memory**
    - Establish a shared memory region, read/write to shared region
    - Accesses are treated as routine memory accesses
    - Faster
  - **Message passing**
    - Exchange message
    - Require kernel intervention
    - Easier to implement in distributed system

# Communications Models



Message passing

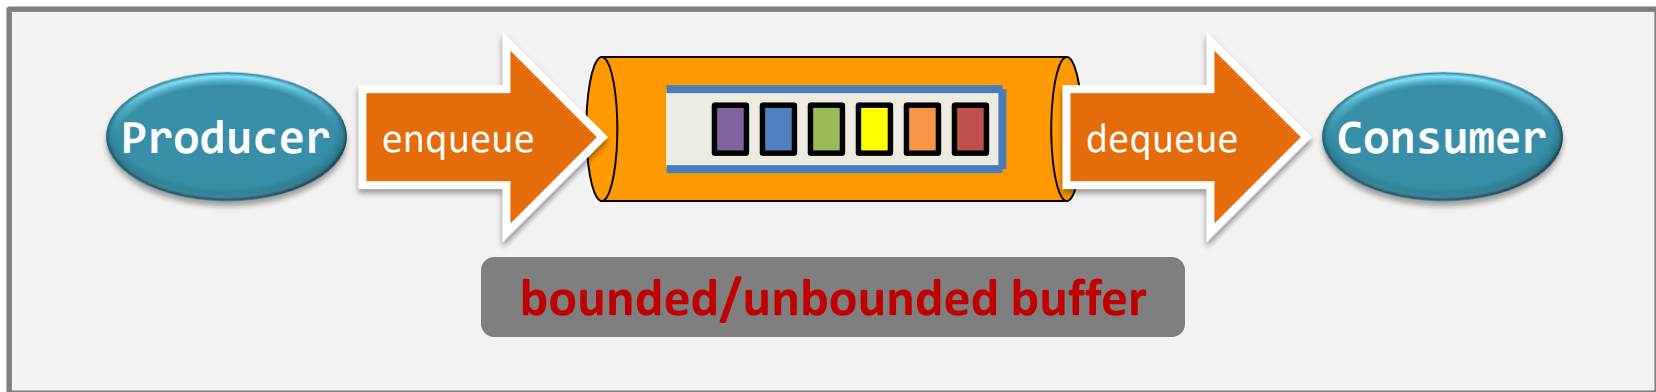


Shared memory



# Producer-Consumer Problem

- Shared memory solution
  - A buffer is needed to allow processes to run concurrently



<b>A buffer</b>	<ul style="list-style-type: none"><li>-It is a shared object;</li><li>-It is a queue (imagine that it is an array implementation of queue).</li></ul>
<b>A producer process</b>	<ul style="list-style-type: none"><li>-It produces a unit of data, and</li><li>-writes that a piece of data to the tail of the buffer at one time.</li></ul>
<b>A consumer process</b>	<ul style="list-style-type: none"><li>-It removes a unit of data from the head of the bounded buffer at one time.</li></ul>

# Producer-Consumer Problem

- Focus on bounded buffer: what are the requirements?

## Producer-consumer requirement #1

When the producer wants to  
(a) put a new item in the buffer, but  
(b) **the buffer is already full...**

Then,

- (1) **The producer should be suspended**, and
- (2) **The consumer should wake the producer up** after she has dequeued an item.

## Producer-consumer requirement #2

When the consumer wants to  
(a) consumes an item from the buffer, but  
(b) **the buffer is empty...**

Then,

- (1) **The consumer should be suspended**, and
- (2) **The producer should wake the consumer up** after she has enqueued an item.

# Producer-consumer solution (shared mem)

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

Shared memory by producer  
& consumer processes



out (consumer)    in (producer)

**Only allows BUFFER\_SIZE-1  
items at the same time. Why?**

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Producer

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item */
}
```

Consumer

# Message Passing

- Communicating processes may reside on different computers connected by a network
- IPC facility provides two operations:
  - **send**(*message*)
  - **receive**(*message*)

# Message Passing (Cont.)

- If processes  $P$  and  $Q$  wish to communicate
  - Establish a **communication link** between them
  - Exchange messages via send/receive



- Implementation issues (logical):
  - Naming: Direct/indirect communication
  - Synchronization: Synchronous/asynchronous
  - Buffering

# Naming

- How to refer to each other?
- **Direct communication**: explicitly name each other
  - Operations (symmetry)
    - **send** ( $Q$ , *message*) – send a message to process  $Q$
    - **receive**( $P$ , *message*) – receive a message from process  $P$
  - Properties of communication link
    - Links are established automatically (every pair can establish)
    - A link is associated with exactly one pair of processes
    - Between each pair, there exists exactly one link
  - Disadvantage: limited modularity (hard-coding)


# Naming

- How to refer to each other?
- **Indirect communication**: sent to and received from mailboxes (ports)
  - Operations
    - **send** (*A, message*) – send a message to mailbox A
    - **receive**(*A, message*) – receive a message from mailbox A
  - Properties of communication link
    - A link is established between a pair of processes only if both members have a shared mailbox
    - A link may be associated with more than two processes
    - Between each pair, a number of different links may exist

# Issues of Indirect Communication

- ISSUE1: Who receives the message when multiple processes are associated with one link?

- Who gets the message?



- Policies

- Allow a link to be associated with at most two processes
    - Allow only one process at a time to execute a receive operation
    - Allow the system to select arbitrarily the receiver (based on an algorithm). Sender is notified who the receiver was.

- ISSUE2: Who owns the mailbox?

- The process (ownership may be passed)
  - The OS (need a method to create, send/receive, delete)



# Synchronization


- How to implement send/receive?
  - **Blocking** is considered **synchronous**
    - **Blocking send** - the sender is blocked until the msg is received
    - **Blocking receive** - the receiver is blocked until a msg is available
  - **Non-blocking** is considered **asynchronous**
    - **Non-blocking send** - the sender sends the message and resumes
    - **Non-blocking receive** - the receiver receives a valid msg or null
- Different combinations are possible
  - When both send and receive are blocking, we have a *rendezvous* between the processes.
  - Other combinations need *buffering*.

# Buffering

- Messages reside in a temporary queue, which can be implemented in three ways
  - **Zero capacity** – no messages are queued on a link, sender must wait for receiver (no buffering)
  - **Bounded capacity** – finite length of  $n$  messages, sender must wait if link is full
  - **Unbounded capacity** – infinite length, sender never waits

# Inter-process communication (IPC)

- What and how?
- POSIX shared memory



# POSIX Shared Memory

- POSIX shared memory is organized using memory-mapped file
  - Associate the region of shared memory with a file
- Illustrate with the producer-consumer problem
  - Producer
  - Consumer

# POSIX Shared Memory

- Producer

- Create a shared-memory object

- `shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666) ;`

Name of the shared memory object

Create the object if it does not exist

Open for reading & writing

Directory permissions

# POSIX Shared Memory

- Producer
  - Create a shared-memory object
    - `shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
  - Configure object size
    - `ftruncate(shm_fd, SIZE);`

File descriptor for the shared mem. Obj.

Size of the shared-memory object

# POSIX Shared Memory

- Producer
  - Create a shared-memory object
    - `shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
  - Configure object size
    - `ftruncate(shm_fd, SIZE);`
  - Establish a memory-mapped file containing the object
    - `ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);`
      - Allows writing to the object  
(only writing is necessary for producer)
      - Changes to the shared-memory object will  
be visible to all processes sharing the object

# POSIX Shared Memory

- Consumer

- Open the shared-memory object

- `shm_fd = shm_open(name, O_RDONLY, 0666) ;`

Open for read only



# POSIX Shared Memory

- Consumer
  - Open the shared-memory object
    - `shm_fd = shm_open(name, O_RDONLY, 0666);`
  - Memory map the object
    - `ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);`

Allows reading to the object  
(only reading is necessary for consumer)

# POSIX Shared Memory

- Consumer
  - Open the shared-memory object
    - `shm_fd = shm_open(name, O_RDONLY, 0666);`
  - Memory map the object
    - `ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);`
  - Remove the shared memory object
    - `shm_unlink(name);`

# POSIX Shared Memory – Complete Solution

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
```

## Producer

```
int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
```

## Consumer

```
int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);


    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

**Direct access to the shared memory region**

# Inter-process communication (IPC)

- What and how?
- POSIX shared memory
- Sockets




# Sockets

- A **socket** is defined as an endpoint for communication (over a network)
  - A pair of processes employ a pair of sockets
  - A socket is identified by an **IP address** and a **port** number
  - All ports below 1024 are used for standard services
    - telnet server listens to port 23
    - FTP server listens to port 21
    - HTTP server listens to port 80

# Sockets

- Socket uses a client-server architecture
  - Server waits for incoming client requests by listening to a specific port
  - Accepts a connection from the client socket to complete the connection
- All connections must be unique
  - Establishing a new connection on the same host needs another port (>1024)
- Special IP address 127.0.0.1 (**loopback**) refers to itself
  - Allow a client and server on the same host to communicate using the TCP/IP protocol



# Example in Java

- Three types of sockets
  - **Connection-oriented (TCP)**, **Connectionless (UDP)**, **Multicast** – data can be sent to multiple recipients

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);
            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);


            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

# Inter-process communication (IPC)


- What and how?
- POSIX shared memory
- Sockets
- Pipes





# What is pipe?

- Pipe is a **shared object**.
  - Using pipe is a way to realize IPC.
  - Acts as a conduit allowing two processes to communicate.



# Pipes

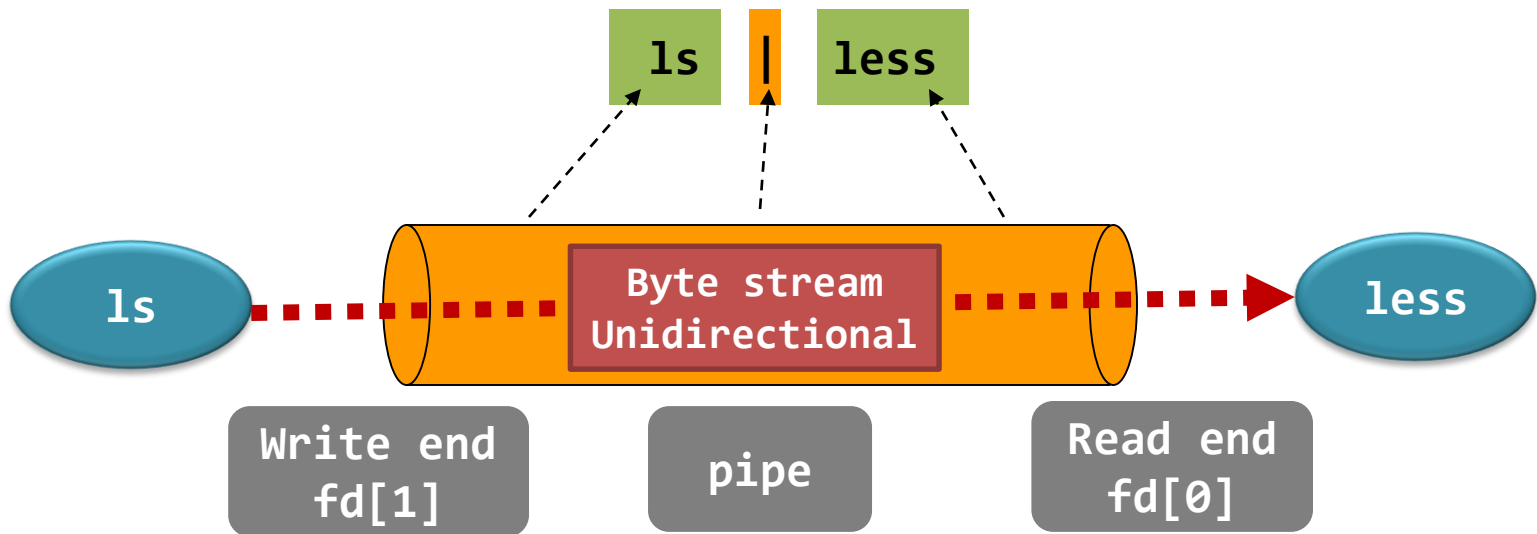
- Four issues:
  - Is the communication unidirectional or bidirectional?
  - In the case of two-way communication, is it half or full-duplex?
  - Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?
  - Can the pipes be used over a network?
- Two common pipes
  - Ordinary pipes and named pipes

# Ordinary Pipes

- Ordinary pipes (no name in file system)
  - Ordinary pipes are used only for related processes (**parent-child relationship**)
    - Processes must reside on the same machine
  - Ordinary pipes are **unidirectional** (one-way communication)
  - Ceases to exist after communication has finished
- Ordinary pipes allow communication in standard producer-consumer style
  - Producer writes to one end (**write-end**)
  - Consumer reads from the other end (**read-end**)


# UNIX Pipe

- UNIX treats a pipe as a special file (child inherits it from parent)
  - Create: `pipe(int fd[]);`
    - `fd[0]`: read end
    - `fd[1]`: write end
  - Access: Ordinary `read()` and `write()` system calls



# UNIX Pipe

- Pipes are anonymous (no name in file system), then how to share?
  - **fork()** duplicates parent's file descriptors
  - Parent and child use each end of the pipe



# UNIX Pipe

```
/* fork a child process */  
pid = fork();
```

Create a child process

```
if (pid < 0) { /* error occurred */  
    fprintf(stderr, "Fork Failed");  
    return 1;  
}
```

```
if (pid > 0) { /* parent process */  
    /* close the unused end of the pipe */  
    close(fd[READ_END]);  
  
    /* write to the pipe */  
    write(fd[WRITE_END], write_msg, strlen(write_msg)+1);  
  
    /* close the write end of the pipe */  
    close(fd[WRITE_END]);  
}
```

Parent process  
Use the write end only

```
else { /* child process */  
    /* close the unused end of the pipe */  
    close(fd[WRITE_END]);  
  
    /* read from the pipe */  
    read(fd[READ_END], read_msg, BUFFER_SIZE);  
    printf("read %s", read_msg);  
  
    /* close the read end of the pipe */  
    close(fd[READ_END]);  
}
```

unidirectional (one-  
way communication)

Child process  
Use the read end only

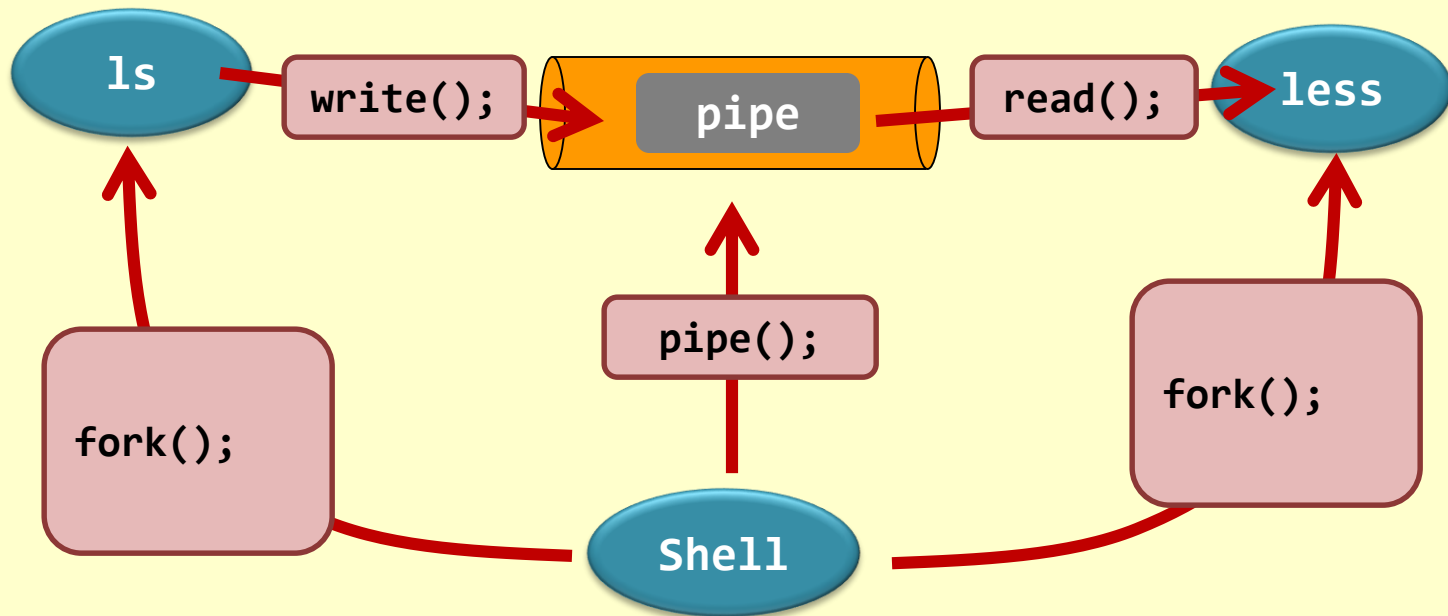
# Pipe - Shell Example

ls

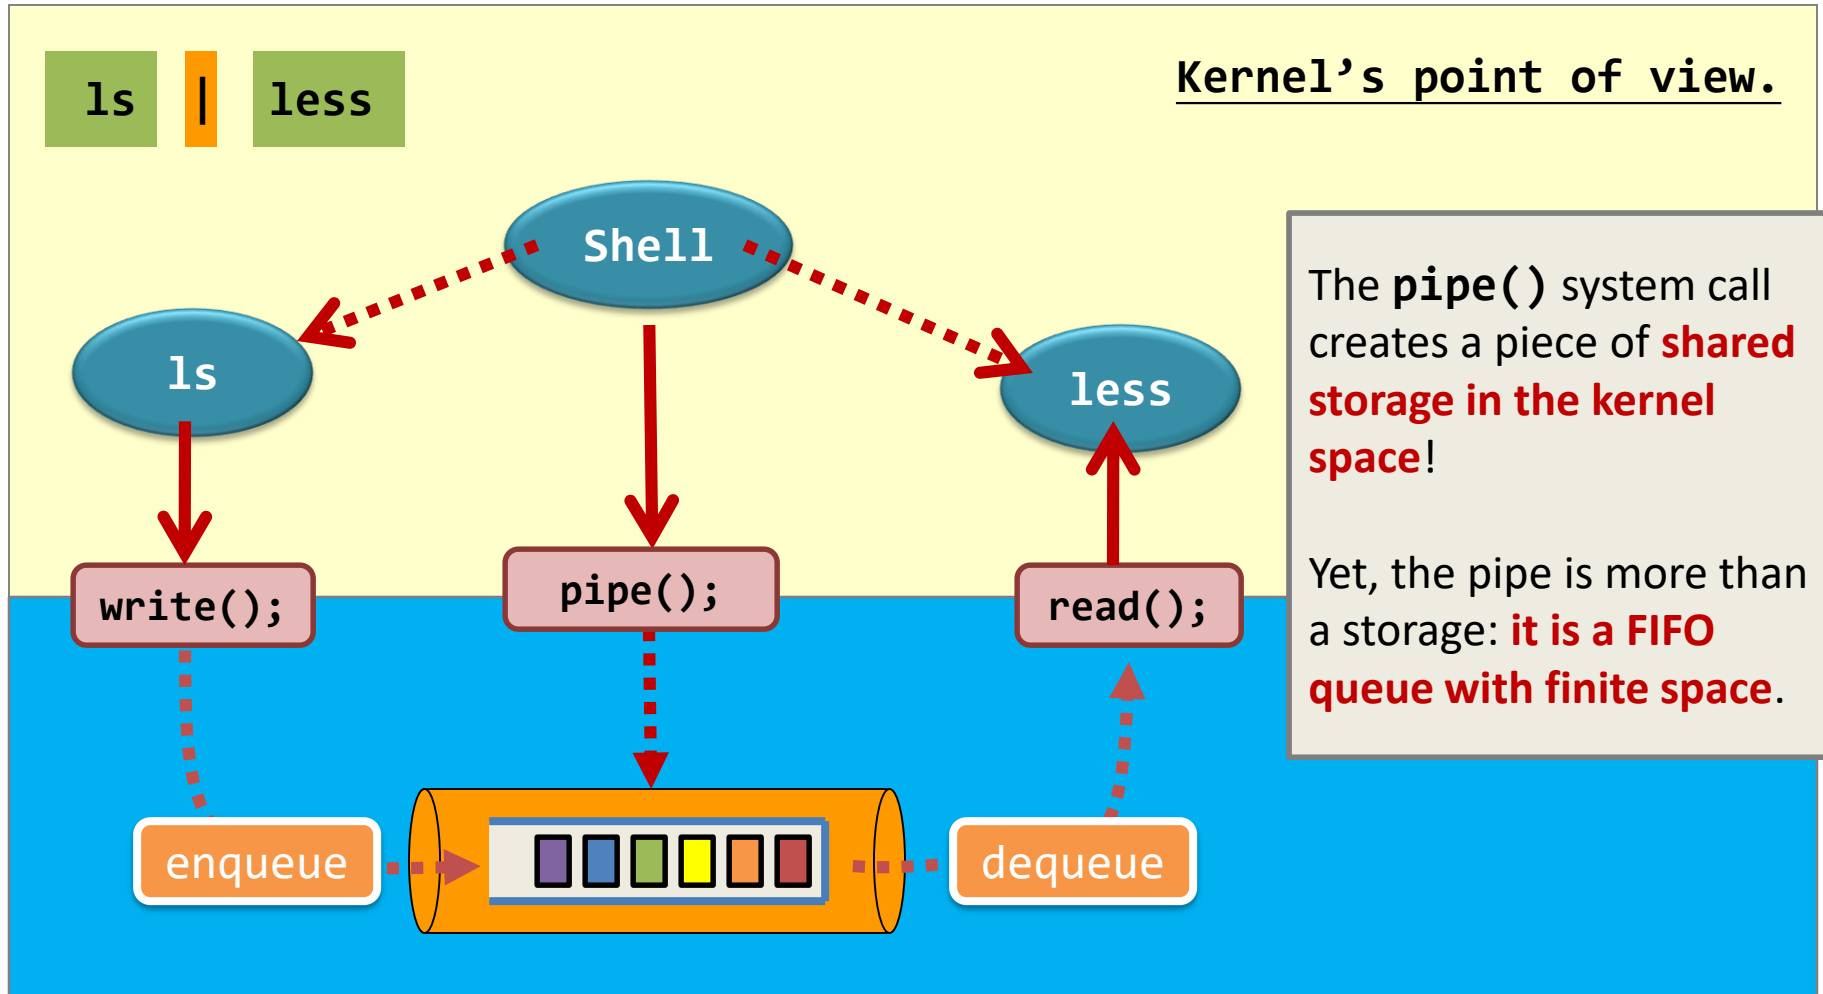
|

less

Programmer's point of view.

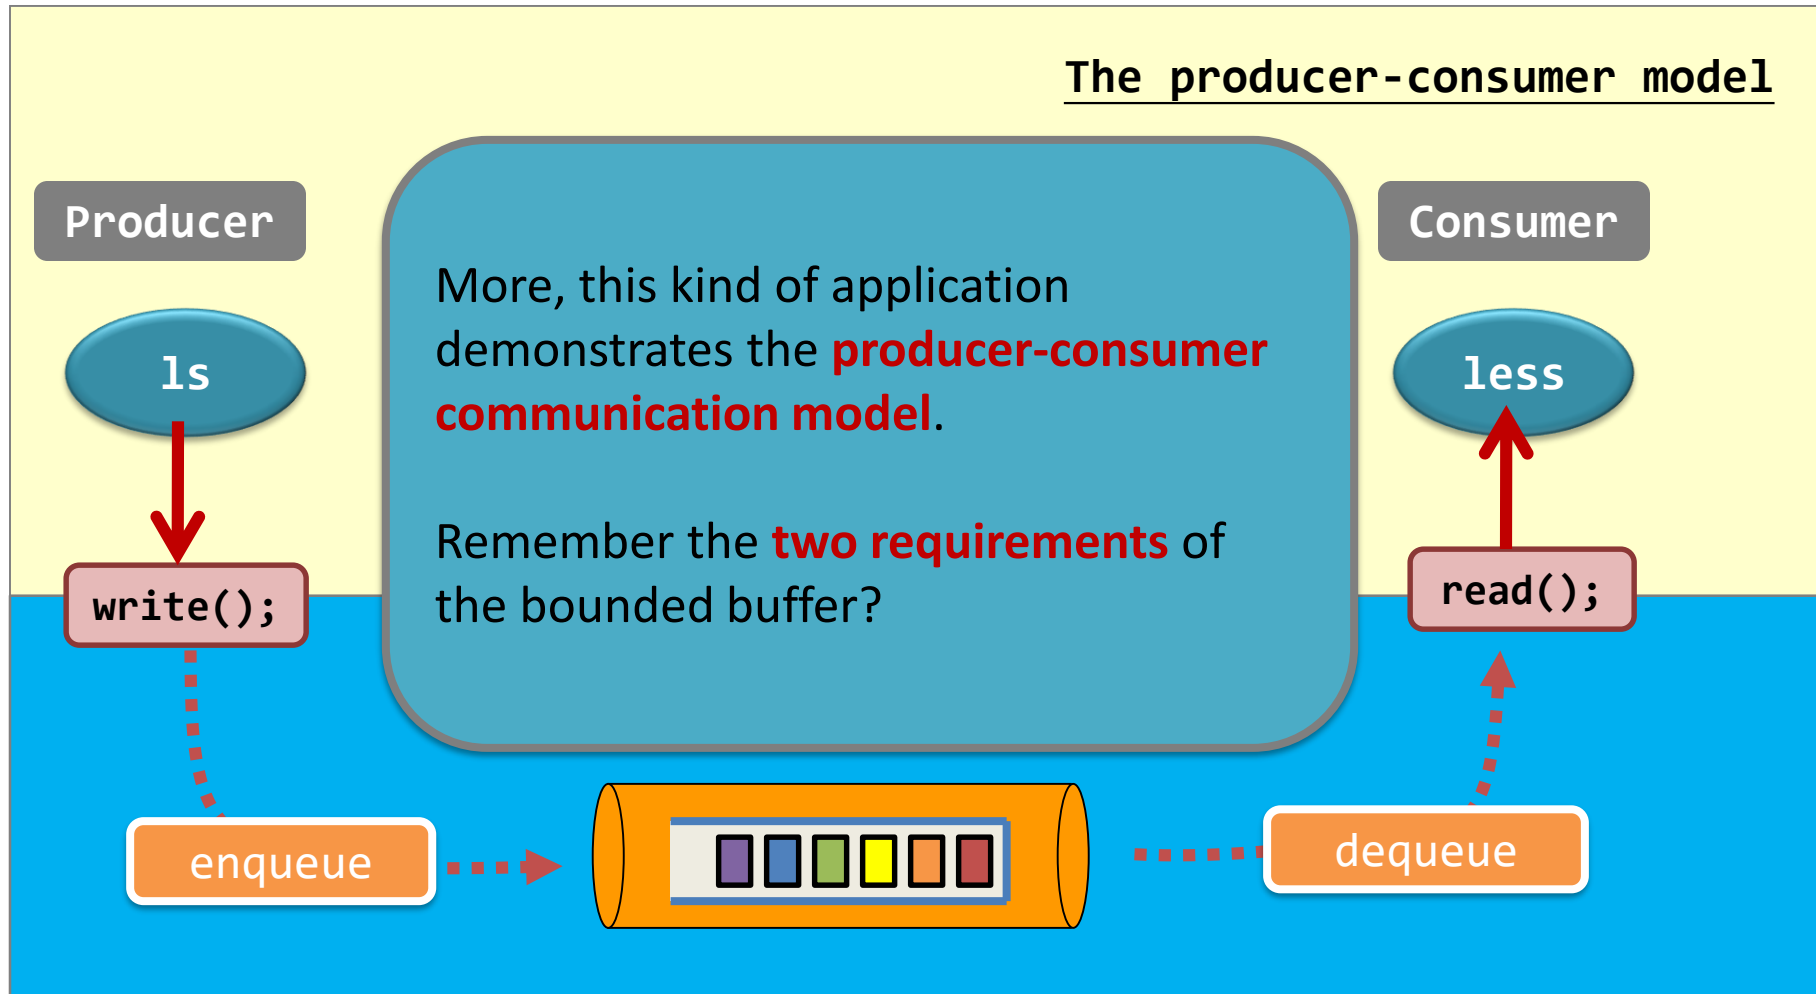


# Pipe – Shell Example





# Pipe – Shell Example



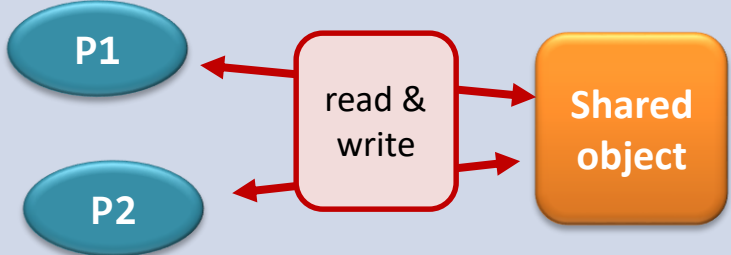
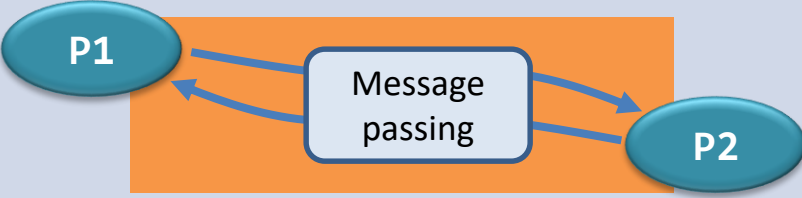
# Named Pipes

- Named pipes (pipe with name in file system)
  - No parent-child relationship is necessary (processes must reside on the same machine)
  - Several processes can use the named pipe for communication (may have several writers)
  - Continue to exist until it is explicitly deleted
  - Communication is bidirectional (still half-duplex)
- Named pipes are referred to as **FIFOs** in UNIX
  - Treated as typical files
  - `mkfifo()`, `open()`, `read()`, `write()`, `close()`

# Story so far...

- Interprocess communication (IPC)
  - Necessary for cooperating processes
  - Producer-consumer model
- IPC models
  - Shared memory & message passing
- IPC schemes
  - Shared memory
  - Ordinary pipes (parent-child processes)
  - FIFOs (processes on the same machine)
  - Sockets (intermachine communication)
- More: Michael Kerrisk, “The Linux Programming Interface” (<http://www.man7.org/tlpi/>)

# IPC models – another point of view

Shared Objects	Message Passing
 <p>The diagram shows two processes, P1 and P2, represented as blue ovals on the left. To their right is a red rounded rectangle labeled 'read &amp; write' with arrows pointing to an orange rounded rectangle labeled 'Shared object'. This indicates that both processes interact with a common shared object through read and write operations.</p>	 <p>The diagram shows two processes, P1 and P2, represented as blue ovals. Between them is a light blue rounded rectangle labeled 'Message passing' set against an orange background. Arrows show communication flow from P1 to the message passing box and from the box to P2, representing a mediated communication channel.</p>
<p><b><u>Challenge.</u></b> Coordination can only be done by detecting the status of the shared object. <i>E.g., is the pipe empty / full?</i></p>	<p><b><u>Challenge.</u></b> Coordination relies on the reliability and the efficiency of the communication medium (and protocol).</p>
<p>E.g., pipes, shared memory, and regular files.</p>	<p>E.g., socket programming, message passing interface (MPI) library.</p>